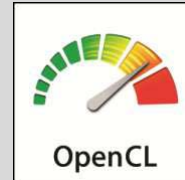


The Open Computing Language (OpenCL)

Mike Bailey

mjb@cs.oregonstate.edu

Oregon State University



Oregon State University
Computer Graphics

openc1.pptx

mjb - February 19, 2015

OpenCL

- Consists of two parts: a C/C++-callable API and a C-ish programming language. It was originally proposed by Apple, but now is a multi-vendor standard
- The programming language can run on NVIDIA GPUs, AMD GPUs, Intel CPUs, Intel GPUs, and (supposedly) FPGAs. But, OpenCL is at its best on compute devices with large amounts of **data parallelism**, which usually implies GPUs.
- You break your computational problem up into small pieces. Each piece gets farmed out to threads on the GPU in a SPMD way.
- OpenCL can share data, and interoperate with, OpenGL
- There is a JavaScript implementation of OpenCL, called WebCL
- There is a JavaScript implementation of OpenGL, called WebGL
- WebCL can share data, and interoperate with, WebGL
- The OpenCL C-ish programming language cannot do recursion, cannot use pointers, and cannot use a stack.



Oregon State University
Computer Graphics

mjb - February 19, 2015

Who Is Behind OpenCL? Members of Khronos's OpenCL Working Group



Oregon State University
Computer Graphics

mjb - February 19, 2015

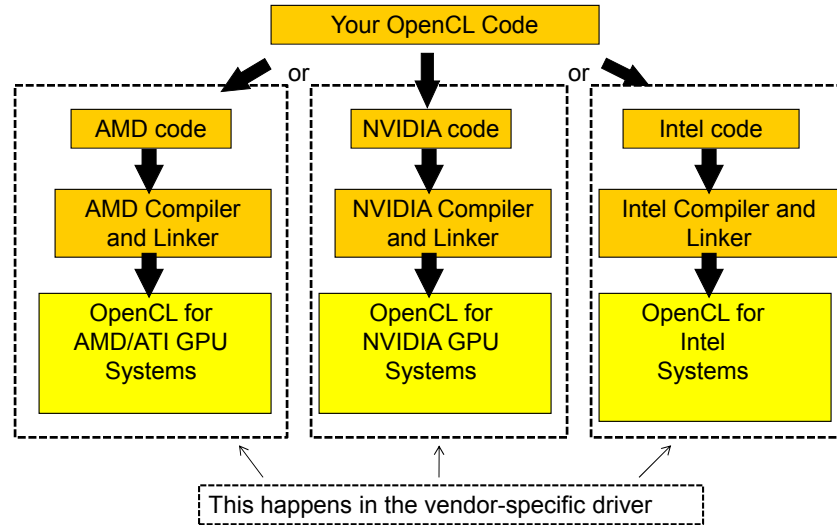
Example of using OpenCL in a System-on-a-Chip: Qualcomm Node - Full Linux and OpenCL



Oregon State University
Computer Graphics

mjb - February 19, 2015

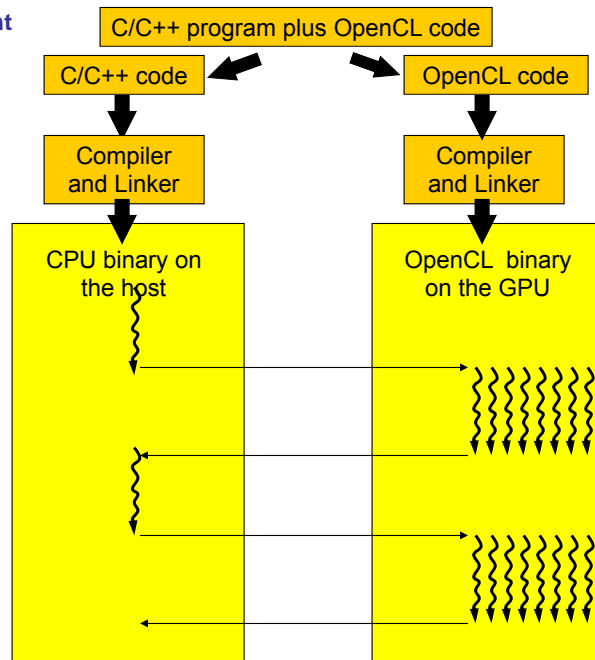
OpenCL – Vendor-independent GPU Programming



Oregon State University
Computer Graphics

mjb – February 19, 2015

The OpenCL Programming Environment



Oregon State University
Computer Graphics

mjb – February 19, 2015

OpenCL wants you to break the problem up into Pieces

If you were writing in C/C++, you would say:

```
void
ArrayMult( int n, float *a, float *b, float *c)
{
    for ( int i = 0; i < n; i++ )
        c[i] = a[i] * b[i];
}
```

If you were writing in OpenCL, you would say:

```
kernel
void
ArrayMult( global float *dA, global float *dB, global float *dC)
{
    int gid = get_global_id ( 0 );
    dC[gid] = dA[gid] * dB[gid];
}
```

This is basically PCAM with lots of P and little A



Oregon State University
Computer Graphics

mjb - February 19, 2015

The OpenCL Language also supports Vector Parallelism

OpenCL code can be vector-oriented, meaning that it can perform a single instruction on multiple data values at the same time (SIMD).

Vector data types are: `char n` , `int n` , `float n` , where $n = 2, 4, 8, \text{ or } 16$.

```
float4 f, g;
f = (float4)( 1.f, 2.f, 3.f, 4.f );

float16 a16, x16, y16, z16;

f.x = 0.;
f.xy = g.zw;
x16.s89ab = f;

float16 a16 = x16 * y16 + z16;
```

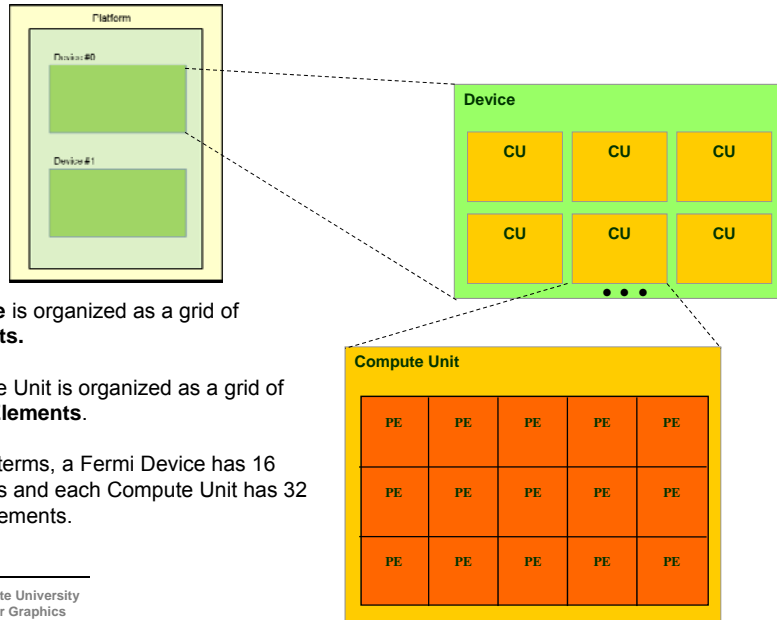
(Note: just because the language supports it, doesn't mean the hardware does.)



Oregon State University
Computer Graphics

mjb - February 19, 2015

From the GPU101 Notes:
Compute Units and Processing Elements are Arranged in Grids



A GPU **Device** is organized as a grid of **Compute Units**.

Each Compute Unit is organized as a grid of **Processing Elements**.

So in NVIDIA terms, a Fermi Device has 16 Compute Units and each Compute Unit has 32 Processing Elements.



Oregon State University
Computer Graphics

mjo - February 19, 2015

OpenCL Software Terminology:
Work-Groups and Work-Items are Arranged in Grids

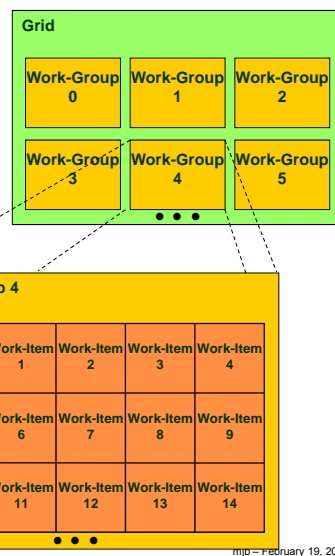
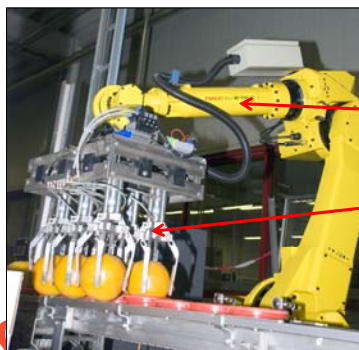
An OpenCL program is organized as a grid of **Work-Groups**.

Each Work-Group is organized as a grid of **Work-Items**.

In terms of hardware, a Work-Group runs on a Compute Unit and a Work-Item runs on a Processing Element (PE).

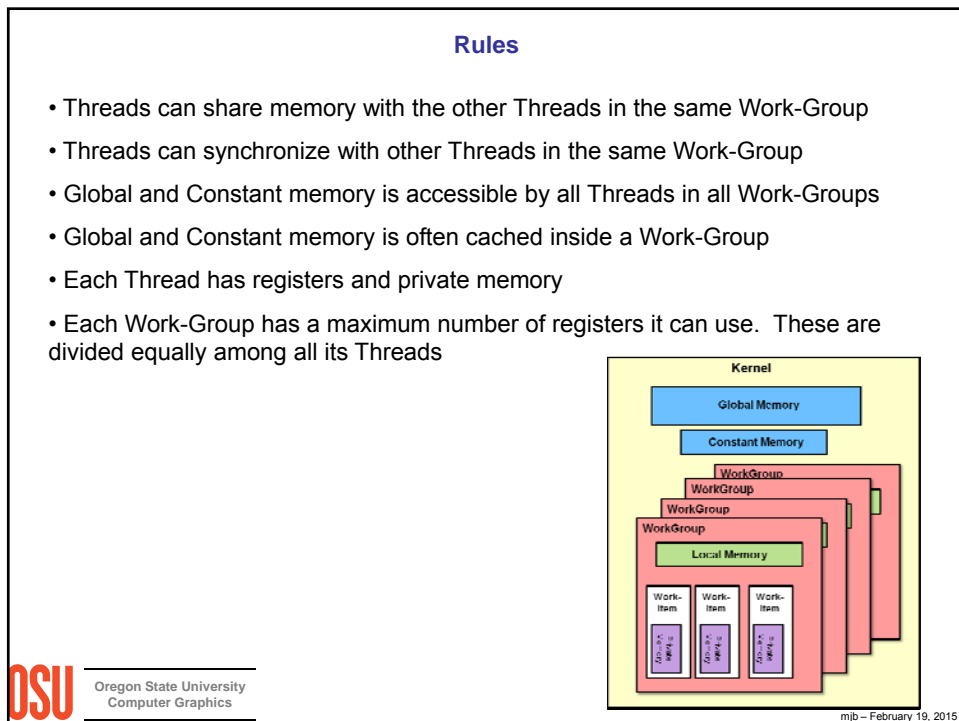
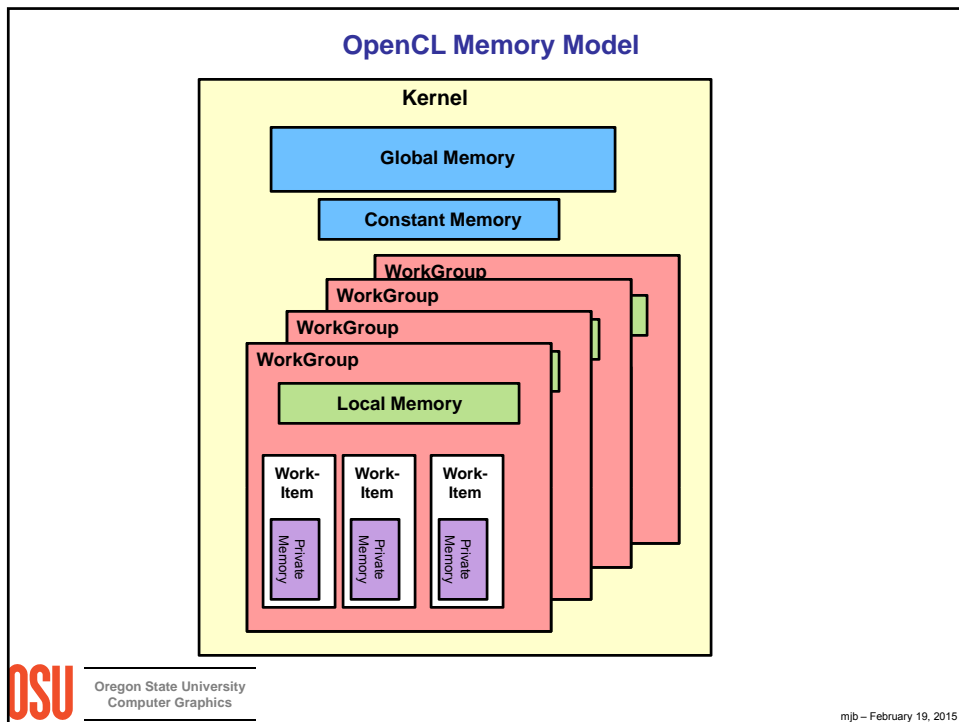
One thread is assigned to each Work-Item.

Threads are swapped on and off the PEs.



Computer Graphics <http://news.cision.com>

mjo - February 19, 2015

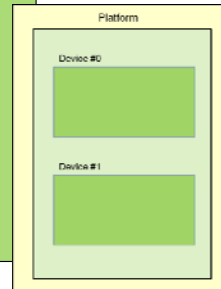


Querying the Number of Platforms (usually one)

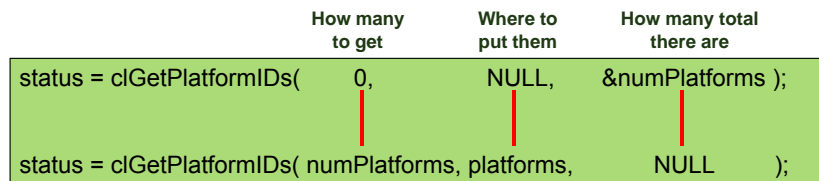
```
cl_uint numPlatforms;
status = clGetPlatformIDs( 0, NULL, &numPlatforms );
if( status != CL_SUCCESS )
    fprintf( stderr, "clGetPlatformIDs failed (1)\n" );

fprintf( stderr, "Number of Platforms = %d\n", numPlatforms );

cl_platform_id * platforms = new cl_platform_id[ numPlatforms ];
status = clGetPlatformIDs( numPlatforms, platforms, NULL );
if( status != CL_SUCCESS )
    fprintf( stderr, "clGetPlatformIDs failed (2)\n" );
```



This way of querying information is a recurring OpenCL pattern:



Oregon State University
Computer Graphics

mjb - February 19, 2015

OpenCL Error Codes

This one is #define'd as zero.
All the others are negative.

CL_SUCCESS	CL_INVALID_QUEUE_PROPERTIES
CL_DEVICE_NOT_FOUND	CL_INVALID_COMMAND_QUEUE
CL_DEVICE_NOT_AVAILABLE	CL_INVALID_HOST_PTR
CL_COMPILER_NOT_AVAILABLE	CL_INVALID_MEM_OBJECT
CL_MEM_OBJECT_ALLOCATION_FAILURE	CL_INVALID_IMAGE_FORMAT_DESCRIPTOR
CL_OUT_OF_RESOURCES	CL_INVALID_IMAGE_SIZE
CL_OUT_OF_HOST_MEMORY	CL_INVALID_SAMPLER
CL_PROFILING_INFO_NOT_AVAILABLE	CL_INVALID_BINARY
CL_MEM_COPY_OVERLAP	CL_INVALID_BUILD_OPTIONS
CL_IMAGE_FORMAT_MISMATCH	CL_INVALID_PROGRAM
CL_IMAGE_FORMAT_NOT_SUPPORTED	CL_INVALID_PROGRAM_EXECUTABLE
CL_BUILD_PROGRAM_FAILURE	CL_INVALID_KERNEL_NAME
CL_MAP_FAILURE	CL_INVALID_KERNEL_DEFINITION
CL_INVALID_VALUE	CL_INVALID_KERNEL
CL_INVALID_DEVICE_TYPE	CL_INVALID_ARG_INDEX
CL_INVALID_PLATFORM	CL_INVALID_ARG_VALUE
CL_INVALID_DEVICE	CL_INVALID_ARG_SIZE
CL_INVALID_CONTEXT	CL_INVALID_KERNEL_ARGS
	CL_INVALID_WORK_DIMENSION



Oregon State University
Computer Graphics

mjb - February 19, 2015

A Way to Print OpenCL Error Codes – get from the Class Web Site

```

struct errorcode
{
    cl_int      statusCode;
    char *      meaning;
};

ErrorCodes[ ] =
{
    { CL_SUCCESS,          "" },
    { CL_DEVICE_NOT_FOUND, "Device Not Found" },
    { CL_DEVICE_NOT_AVAILABLE, "Device Not Available" },
    { CL_INVALID_MIP_LEVEL, "Invalid MIP Level" },
    { CL_INVALID_GLOBAL_WORK_SIZE, "Invalid Global Work Size" },
};

void
PrintCLError( cl_int errorCode, char * prefix, FILE *fp )
{
    if( errorCode == CL_SUCCESS )
        return;

    const int numErrorCodes = sizeof( ErrorCodes ) / sizeof( struct errorcode );
    char * meaning = " ";
    for( int i = 0; i < numErrorCodes; i++ )
    {
        if( errorCode == ErrorCodes[i].statusCode )
        {
            meaning = ErrorCodes[i].meaning;
            break;
        }
    }

    fprintf( fp, "%s %s\n", prefix, meaning );
}
    
```



Oreg
Co

mjb – February 19, 2015

Querying the Number of Devices on a Platform

```

// find out how many devices are attached to each platform and get their ids:

status = clGetDeviceIDs( platform, CL_DEVICE_TYPE_ALL, 0, NULL, &numDevices );

devices = new cl_device_id[ numDevices ];

status = clGetDeviceIDs( platform, CL_DEVICE_TYPE_ALL, numDevices, devices, NULL );
    
```

Getting Just the GPU Device

```

cl_device_id device;

status = clGetDeviceIDs( platform, CL_DEVICE_TYPE_GPU, 1, &device, NULL );
    
```



Oregon State University
Computer Graphics

mjb – February 19, 2015

Querying the Device (this is *really* useful!)

```
// find out how many platforms are attached here and get their ids:

cl_uint numPlatforms;
status = clGetPlatformIDs( 0, NULL, &numPlatforms );
if( status != CL_SUCCESS )
    fprintf( stderr, "clGetPlatformIDs failed (1)\n" );

fprintf( OUTPUT, "Number of Platforms = %d\n", numPlatforms );

cl_platform_id *platforms = new cl_platform_id[ numPlatforms ];
status = clGetPlatformIDs( numPlatforms, platforms, NULL );
if( status != CL_SUCCESS )
    fprintf( stderr, "clGetPlatformIDs failed (2)\n" );

cl_uint numDevices;
cl_device_id *devices;

for( int i = 0; i < (int)numPlatforms; i++ )
{
    fprintf( OUTPUT, "Platform #%d:\n", i );
    size_t size;
    char *str;

    clGetPlatformInfo( platforms[i], CL_PLATFORM_NAME, 0, NULL, &size );
    str = new char [ size ];
    clGetPlatformInfo( platforms[i], CL_PLATFORM_NAME, size, str, NULL );
    fprintf( OUTPUT, "\tName = %s\n", str );
    delete[] str;

    clGetPlatformInfo( platforms[i], CL_PLATFORM_VENDOR, 0, NULL, &size );
    str = new char [ size ];
    clGetPlatformInfo( platforms[i], CL_PLATFORM_VENDOR, size, str, NULL );
    fprintf( OUTPUT, "\tVendor = %s\n", str );
    delete[] str;
}
```

February 19, 2015

```
clGetPlatformInfo( platforms[i], CL_PLATFORM_VERSION, 0, NULL, &size );
str = new char [ size ];
clGetPlatformInfo( platforms[i], CL_PLATFORM_VERSION, size, str, NULL );
fprintf( OUTPUT, "\tVersion = %s\n", str );
delete[] str;

clGetPlatformInfo( platforms[i], CL_PLATFORM_PROFILE, 0, NULL, &size );
str = new char [ size ];
clGetPlatformInfo( platforms[i], CL_PLATFORM_PROFILE, size, str, NULL );
fprintf( OUTPUT, "\tProfile = %s\n", str );
delete[] str;

// find out how many devices are attached to each platform and get their ids:


status = clGetDeviceIDs( platforms[i], CL_DEVICE_TYPE_ALL, 0, NULL, &numDevices );
if( status != CL_SUCCESS )
    fprintf( stderr, "clGetDeviceIDs failed (2)\n" );

devices = new cl_device_id[ numDevices ];
status = clGetDeviceIDs( platforms[i], CL_DEVICE_TYPE_ALL, numDevices, devices, NULL );
if( status != CL_SUCCESS )
    fprintf( stderr, "clGetDeviceIDs failed (2)\n" );

for( int j = 0; j < (int)numDevices; j++ )
{
    fprintf( OUTPUT, "\tDevice #%d:\n", i );
    size_t size;
    cl_device_type type;
    cl_uint ui;
    size_t sizes[3] = { 0, 0, 0 };

    clGetDeviceInfo( devices[j], CL_DEVICE_TYPE, sizeof(type), &type, NULL );
    fprintf( OUTPUT, "\tType = 0x%04x = ", type );
}
```

mjb - February 19, 2015



```

switch( type )
{
    case CL_DEVICE_TYPE_CPU:
        fprintf( OUTPUT, "CL_DEVICE_TYPE_CPU\n" );
        break;
    case CL_DEVICE_TYPE_GPU:
        fprintf( OUTPUT, "CL_DEVICE_TYPE_GPU\n" );
        break;
    case CL_DEVICE_TYPE_ACCELERATOR:
        fprintf( OUTPUT, "CL_DEVICE_TYPE_ACCELERATOR\n" );
        break;
    default:
        fprintf( OUTPUT, "Other...\n" );
        break;
}

clGetDeviceInfo( devices[i], CL_DEVICE_VENDOR_ID, sizeof(ui), &ui, NULL );
fprintf( OUTPUT, "\t\tDevice Vendor ID = 0x%04x\n", ui );

clGetDeviceInfo( devices[i], CL_DEVICE_MAX_COMPUTE_UNITS, sizeof(ui), &ui, NULL );
fprintf( OUTPUT, "\t\tDevice Maximum Compute Units = %d\n", ui );

clGetDeviceInfo( devices[i], CL_DEVICE_MAX_WORK_ITEM_DIMENSIONS, sizeof(ui), &ui, NULL );
fprintf( OUTPUT, "\t\tDevice Maximum Work Item Dimensions = %d\n", ui );

clGetDeviceInfo( devices[i], CL_DEVICE_MAX_WORK_ITEM_SIZES, sizeof(sizes), sizes, NULL );
fprintf( OUTPUT, "\t\tDevice Maximum Work Item Sizes = %d x %d x %d\n", sizes[0], sizes[1], sizes[2] );

clGetDeviceInfo( devices[i], CL_DEVICE_MAX_WORK_GROUP_SIZE, sizeof(size), &size, NULL );
fprintf( OUTPUT, "\t\tDevice Maximum Work Group Size = %d\n", size );


clGetDeviceInfo( devices[i], CL_DEVICE_MAX_CLOCK_FREQUENCY, sizeof(ui), &ui, NULL );
fprintf( OUTPUT, "\t\tDevice Maximum Clock Frequency = %d MHz\n", ui );
}
        
```

mjb - February 19, 2015

Typical Values from Querying the Device

```

Number of Platforms = 1
Platform #0:
    Name = 'NVIDIA CUDA'
    Vendor = 'NVIDIA Corporation'
    Version = 'OpenCL 1.1 CUDA 4.1.1'
    Profile = 'FULL_PROFILE'
    Device #0:
        Type = 0x0004 = CL_DEVICE_TYPE_GPU
        Device Vendor ID = 0x10de
        Device Maximum Compute Units = 15
        Device Maximum Work Item Dimensions = 3
        Device Maximum Work Item Sizes = 1024 x 1024 x 64
        Device Maximum Work Group Size = 1024
        Device Maximum Clock Frequency = 1401 MHz
        Kernel Maximum Work Group Size = 1024
        Kernel Compile Work Group Size = 0 x 0 x 0
        Kernel Local Memory Size = 0
        
```



Oregon State University
Computer Graphics

mjb - February 19, 2015

Querying to see what extensions are supported on this device

```
size_t extensionSize;

clGetDeviceInfo( device, CL_DEVICE_EXTENSIONS, 0, NULL, &extensionSize );
char *extensions = new char [extensionSize];
clGetDeviceInfo( devices, CL_DEVICE_EXTENSIONS, extensionSize, extensions, NULL );

fprintf( stderr, "\nDevice Extensions:\n" );
for( int i = 0; i < (int)strlen(extensions); i++ )
{
    if( extensions[ i ] == ' ' )
        extensions[ i ] = '\n';
}
fprintf( stderr, "%s\n", extensions );
delete [ ] extensions;
```



Oregon State University
Computer Graphics

mjb - February 19, 2015

Querying to see what extensions are supported on this device

This is the big one
you are looking for

Device Extensions:

cl_khr_byte_addressable_store
cl_khr_icd
cl_khr_gl_sharing
cl_nv_d3d9_sharing
cl_nv_d3d10_sharing
cl_khr_d3d10_sharing
cl_nv_d3d11_sharing
cl_nv_compiler_options
cl_nv_device_attribute_query
cl_nv_pragma_unroll

This one is handy too

cl_khr_global_int32_base_atomics
cl_khr_global_int32_extended_atomics
cl_khr_local_int32_base_atomics
cl_khr_local_int32_extended_atomics
cl_khr_fp64



Oregon State University
Computer Graphics

mjb - February 19, 2015

Steps in Creating and Running an OpenCL program

1. Program header
2. Allocate the host memory buffers
3. Create an OpenCL context
4. Create an OpenCL command queue
5. Allocate the device memory buffers
6. Write the data from the host buffers to the device buffers
7. Read the kernel code from a file
8. Compile and link the kernel code
9. Create the kernel object
10. Setup the arguments to the kernel object
11. Enqueue the kernel object for execution
12. Read the results buffer back from the device to the host
13. Clean everything up



Oregon State University
Computer Graphics

mjb - February 19, 2015

1. .cpp Program Header

```
#include <stdio.h>
#include <math.h>
#include <string.h>
#include <stdlib.h>
#include <omp.h>    // for timing

#include "cl.h"
```



Oregon State University
Computer Graphics

mjb - February 19, 2015

2. Allocate the Host Memory Buffers

```
// allocate the host memory buffers:
float * hA = new float [ NUM_ELEMENTS ];
float * hB = new float [ NUM_ELEMENTS ];
float * hC = new float [ NUM_ELEMENTS ];

// fill the host memory buffers:
for( int i = 0; i < NUM_ELEMENTS; i++ )
{
    hA[ i ] = hB[ i ] = sqrtf( (float) i );
}

// array size in bytes (will need this later):
size_t dataSize = NUM_ELEMENTS * sizeof( float );

// opengl function return status:
cl_int status; // test against CL_SUCCESS
```

This could have also been done like this:

```
float hA[ NUM_ELEMENTS ];
```

Global memory and the heap typically have more space than the stack. So, rarely do you want to allocate large arrays like this as a local variable.

(Here, it's being done on the heap.)



Oregon State University
Computer Graphics

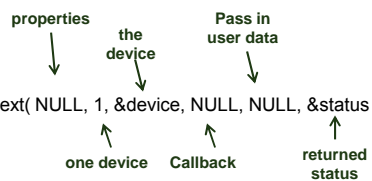
mjb - February 19, 2015

3. Create an OpenCL Context

```
cl_context context = clCreateContext( NULL, 1, &device, NULL, NULL, &status );
```

// create a context:

```
cl_context context = clCreateContext( NULL, 1, &device, NULL, NULL, &status );
```



Oregon State University
Computer Graphics

mjb - February 19, 2015

6. Write the Data from the Host Buffers to the Device Buffers

// enqueue the 2 commands to write data into the device buffers:

```
status = clEnqueueWriteBuffer( cmdQueue, dA, CL_FALSE, 0, dataSize, hA, 0, NULL, NULL );  
status = clEnqueueWriteBuffer( cmdQueue, dB, CL_FALSE, 0, dataSize, hB, 0, NULL, NULL );
```

command queue want to block until done? # bytes # events event object

status = clEnqueueWriteBuffer(cmdQueue, dA, CL_FALSE, 0, dataSize, hA, 0, NULL, NULL);

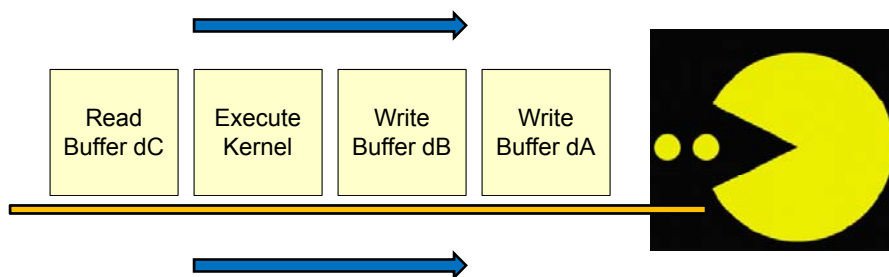
device buffer offset host buffer event wait list



Oregon State University
Computer Graphics

mjb - February 19, 2015

Enqueuing Works Like a Conveyor Belt

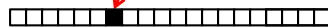


Oregon State University
Computer Graphics

mjb - February 19, 2015

The .cl File

```
kernel
void
ArrayMult( global const float *dA, global const float *dB, global float *dC )
{
    int gid = get_global_id( 0 );
    dC[gid] = dA[gid] * dB[gid];
}
```



gid = which element we are dealing with right now.

Which dimension's index are we fetching?

0 = X, 1 = Y, 2 = Z

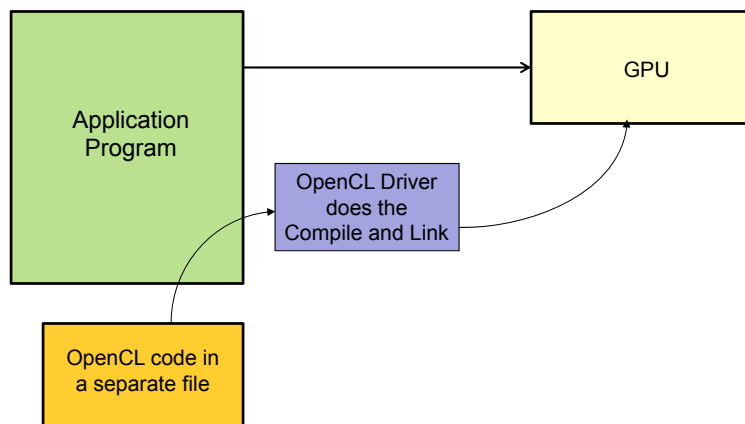
Since this is a 1D problem, X is the only index we need to get.



Oregon State University
Computer Graphics

mjb - February 19, 2015

OpenCL code is compiled in the Driver . . .



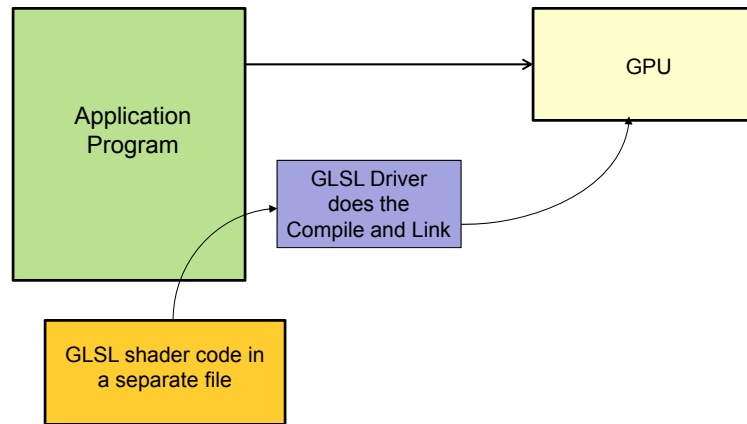
```
kernel void
ArrayMult( global float *A, global float *B, global float *C )
{
    int gid = get_global_id ( 0 );
    C[gid] = A[gid] * B[gid];
}
```



Computer Graphics

mjb - February 19, 2015

(... just like OpenGL's GLSL Shader code is compiled in the driver)



```
void main( )
{
    vec3 newcolor = texture2D( uTexUnit, vST ).rgb;
    newcolor = mix( newcolor, vColor.rgb, uBlend );
    gl_FragColor = vec4(u LightIntensity*newcolor, 1. );
}
```



Oregon State University
Computer Graphics

mjb - February 19, 2015

7. Read the Kernel Code from a File into a Character Array

```
const char *CL_FILE_NAME = { "arraymult.cl" };
...
FILE *fp = fopen( CL_FILE_NAME, "r" );
if( fp == NULL )
{
    fprintf( stderr, "Cannot open OpenCL source file '%s'\n", CL_FILE_NAME );
    return 1;
}

// read the characters from the opencl kernel program:

fseek( fp, 0, SEEK_END );
size_t fileSize = ftell( fp );
fseek( fp, 0, SEEK_SET );
char *clProgramText = new char[ fileSize+1 ];
size_t n = fread( clProgramText, 1, fileSize, fp );
clProgramText[fileSize] = '\0';
fclose( fp );
```

"r" should work, since the .cl file is pure ASCII text, but some people report that it doesn't work unless you use "rb"

Watch out for the '\r' + '\n' problem! (See the next slide.)



Oregon State University
Computer Graphics

mjb - February 19, 2015

A Warning about Editing on Windows and Running on Linux

Some of you will end up having strange, unexplainable problems with your csh scripts, .cpp programs, or .cl programs. This could be because you are typing your code in on Windows (using Notepad or Wordpad or Word) and then running it on Linux. Windows likes to insert an extra carriage return ('\r') at the end of each line, which Linux interprets as a garbage character.

You can test this by typing the Linux command:

od -c loop.csh

which will show you all the characters, even the '\r' (which you don't want) and the '\n' (newlines, which you do want).

To get rid of the carriage returns, enter the Linux command:

tr -d '\r' < loop.csh > loop1.csh

Then run loop1.csh

Or, on some systems, there is a utility called *dos2unix* which does this for you:

dos2unix < loop.csh > loop1.csh

Sorry about this. Unfortunately, this is a fact of life when you mix Windows and Linux.



Oregon State University
Computer Graphics

mjb - February 19, 2015

8. Compile and Link the Kernel Code

```
// create the kernel program on the device:

char * strings [ 1 ];           // an array of strings
strings[0] = clProgramText;
cl_program program = clCreateProgramWithSource( context, 1, (const char **)strings, NULL, &status );
delete [ ] clProgramText;

// build the kernel program on the device:

char *options = { "" };
status = clBuildProgram( program, 1, &device, options, NULL, NULL );
if( status != CL_SUCCESS )
{
    // retrieve and print the error messages:
    size_t size;
    clGetProgramBuildInfo( program, devices[0], CL_PROGRAM_BUILD_LOG, 0, NULL, &size );
    cl_char *log = new cl_char[ size ];
    clGetProgramBuildInfo( program, devices[0], CL_PROGRAM_BUILD_LOG, size, log, NULL );
    fprintf( stderr, "clBuildProgram failed:\n%s\n", log );
    delete [ ] log;
}
```



Oregon State University
Computer Graphics

mjb - February 19, 2015

How does that array-of-strings thing actually work?

```
char *ArrayOfStrings[3];
ArrayOfStrings[0] = "...one commonly-used function...";
ArrayOfStrings[1] = "... another commonly-used function. . . ";
ArrayOfStrings[2] = "... the real OpenCL code . . . ";
cl_program program = clCreateProgramWithSource( context, 1, (const char **) ArrayOfStrings, NULL, &status );
```

These are two ways to provide a *single* character buffer:

```
char *buffer[1];
buffer[0] = "... the entire OpenCL code . . . ";
cl_program program = clCreateProgramWithSource( context, 1, (const char **) buffer, NULL, &status );
```

```
char *buffer = "... the entire OpenCL code . . . ";
cl_program program = clCreateProgramWithSource( context, 1, (const char **) &buffer, NULL, &status );
```



Oregon State University
Computer Graphics

mjb - February 19, 2015

Why use an array of strings to hold the OpenCL program, instead of just a single string?

1. You can use the same OpenCL source and insert the appropriate "#defines" at the beginning
2. You can insert a common header file (\approx a .h file)
3. You can simulate a "#include" to re-use common pieces of code



Oregon State University
Computer Graphics

mjb - February 19, 2015

9. Create the Kernel Object

```
cl_kernel kernel = clCreateKernel( program, "ArrayMult", &status );
```



Oregon State University
Computer Graphics

mjb - February 19, 2015

10. Setup the Arguments to the Kernel Object

```
status = clSetKernelArg( kernel, 0, sizeof(cl_mem), &dA );  
status = clSetKernelArg( kernel, 1, sizeof(cl_mem), &dB );  
status = clSetKernelArg( kernel, 2, sizeof(cl_mem), &dB );
```



Oregon State University
Computer Graphics

mjb - February 19, 2015

11. Enqueue the Kernel Object for Execution

```
size_t globalWorkSize[ 3 ] = { NUM_ELEMENT, 1, 1 };
size_t localWorkSize[ 3 ] = { LOCAL_SIZE, 1, 1 };

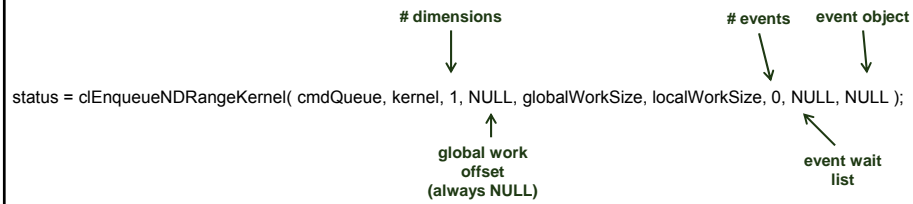
status = clEnqueueBarrier( cmdQueue );

double time0 = omp_get_wtime( );

status = clEnqueueNDRangeKernel( cmdQueue, kernel, 1, NULL, globalWorkSize, localWorkSize, 0, NULL, NULL );

status = clEnqueueBarrier( cmdQueue );

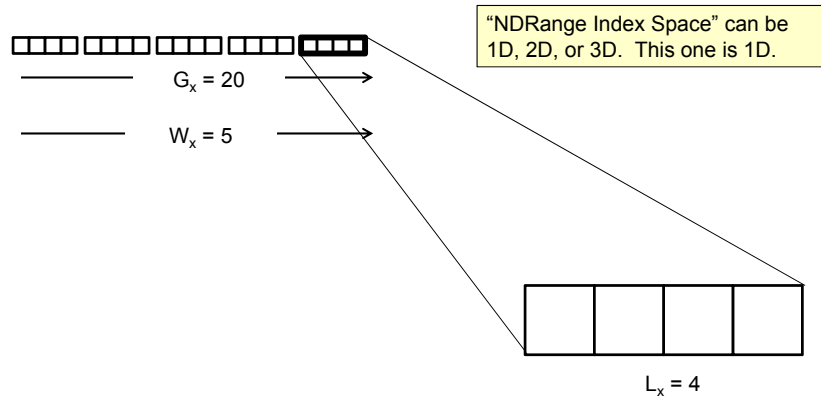
double time1 = omp_get_wtime( );
```



Oregon State University
Computer Graphics

mjb - February 19, 2015

Work-Groups, Local IDs, and Global IDs



$$\#WorkGroups = \frac{GlobalIndexSpaceSize}{WorkGroupSize}$$

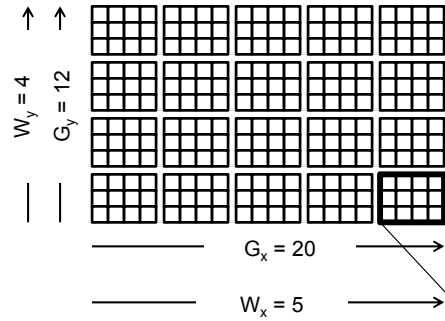
$$5 \times 4 = \frac{20}{4}$$



Oregon State University
Computer Graphics

mjb - February 19, 2015

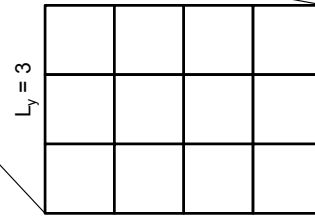
Work-Groups, Local IDs, and Global IDs



"NDRange Index Space" can be 1D, 2D, or 3D. This one is 2D.

$$\#WorkGroups = \frac{GlobalIndexSpaceSize}{WorkGroupSize}$$

$$5 \times 4 = \frac{20 \times 12}{4 \times 3}$$

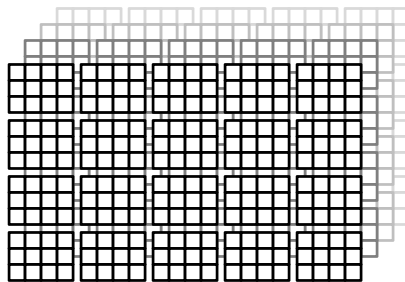


Oregon State University
Computer Graphics

mjb - February 19, 2015

Work-Groups, Local IDs, and Global IDs

"NDRange Index Space" can be 1D, 2D, or 3D. This one is 3D.



Oregon State University
Computer Graphics

mjb - February 19, 2015

Figuring Out What Thread You Are

```

uint    get_work_dim( );

size_t  get_global_size( uint dimindx );

size_t  get_global_id( uint dimindx );

size_t  get_local_size( uint dimindx );

size_t  get_local_id( uint dimindx );

size_t  get_num_groups( uint dimindx );

size_t  get_group_id( uint dimindx );

size_t  get_global_offset( uint dimindx );
    
```

$0 \leq \text{dimindx} \leq 2$

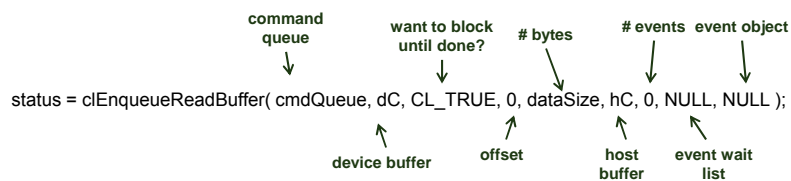


Oregon State University
Computer Graphics

mjb - February 19, 2015

12. Read the Results Buffer Back from the Device to the Host

```
status = clEnqueueReadBuffer( cmdQueue, dC, CL_TRUE, 0, dataSize, hC, 0, NULL, NULL );
```



Oregon State University
Computer Graphics

mjb - February 19, 2015

13. Clean Everything Up

```
// clean everything up:

clReleaseKernel(      kernel );
clReleaseProgram(     program );
clReleaseCommandQueue( cmdQueue );
clReleaseMemObject(   dA );
clReleaseMemObject(   dB );
clReleaseMemObject(   dC );

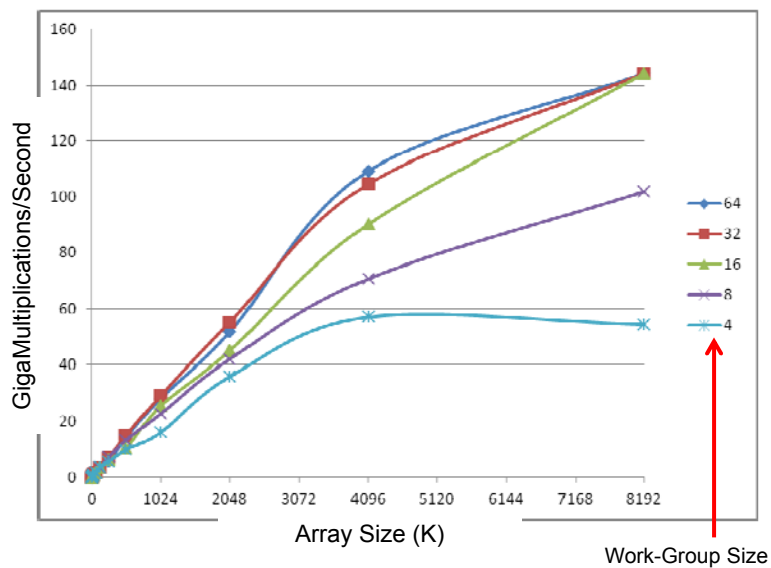
delete [ ] hA;
delete [ ] hB;
delete [ ] hC;
```



Oregon State University
Computer Graphics

mjb - February 19, 2015

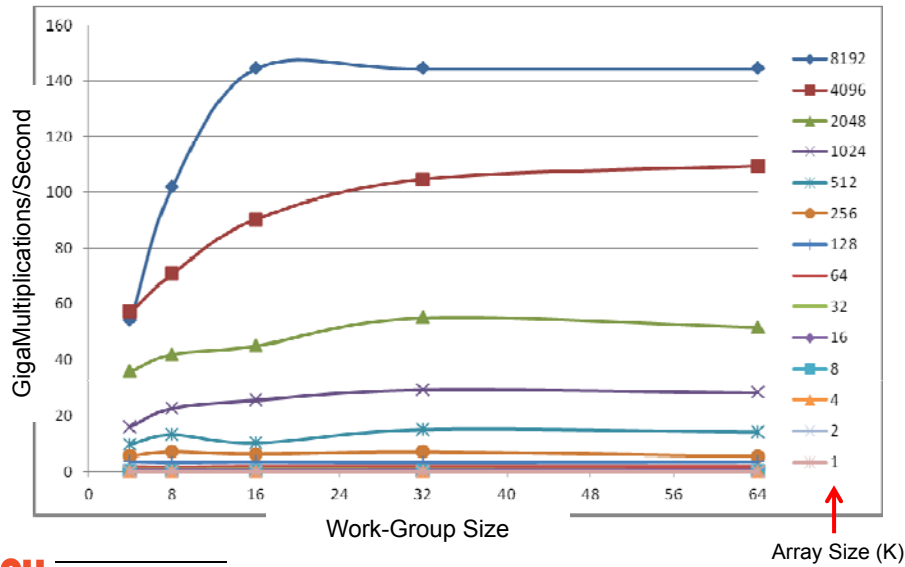
Array Multiplication Performance: What is a Good Work-Group Size?



Oregon State University
Computer Graphics

mjb - February 19, 2015

Array Multiplication Performance: What is a Good Work-Group Size?



Oregon State University
Computer Graphics

mjb - February 19, 2015

Writing the .cl Program's Binary Code

```
size_t binary_sizes;
status = clGetProgramInfo( Program, CL_PROGRAM_BINARY_SIZES, 0, NULL, &binary_sizes );

size_t size;
status = clGetProgramInfo( Program, CL_PROGRAM_BINARY_SIZES, sizeof(size_t), &size, NULL );

unsigned char *binary = new unsigned char [ size ];
status = clGetProgramInfo( Program, CL_PROGRAM_BINARIES, size, &binary, NULL );

FILE *fpbin = fopen( "particles.nv", "wb" );
if( fpbin == NULL )
{
    fprintf( stderr, "Cannot create 'particles.bin'\n" );
}
else
{
    fwrite( binary, 1, size, fpbin );
    fclose( fpbin );
}
delete [ ] binary;
```



Oregon State University
Computer Graphics

mjb - February 19, 2015

Importing that Binary Code back In: 8. Compile and Link the Kernel Code

Instead of doing this:

```
char * strings [ 1 ];
strings[0] = clProgramText;
cl_program program = clCreateProgramWithSource( context, 1, (const char **)strings, NULL, &status );
delete [ ] clProgramText;
```

You would do this:

```
unsigned char byteArray[ numBytes ];
cl_program program = clCreateProgramWithBinary( context, 1, &device, &numBytes, &byteArray, &binaryStatus, &status );
delete [ ] byteArray;
```

And you still have to do this:

```
char *options = { "" };
status = clBuildProgram( program, 1, &device, options, NULL, NULL );
if( status != CL_SUCCESS )
{
    size_t size;
    clGetProgramBuildInfo( program, device, CL_PROGRAM_BUILD_LOG, 0, NULL, &size );
    cl_char *log = new cl_char[ size ];
    clGetProgramBuildInfo( program, device, CL_PROGRAM_BUILD_LOG, size, log, NULL );
    fprintf( stderr, "clBuildProgram failed:\n%s\n", log );
    delete [ ] log;
}
```

osu