

Looking at OpenCL Assembly Language

Mike Bailey

mjb@cs.oregonstate.edu

Oregon State University



opend.assembly.pptx

mjb - March 16, 2015

How to Dump OpenCL Assembly Language

```
size_t size;
status = clGetProgramInfo( Program, CL_PROGRAM_BINARY_SIZES, sizeof(size_t), &size, NULL );
PrintCLError( status, "clGetProgramInfo (1):" );

unsigned char * binary = new unsigned char [ size ];
status = clGetProgramInfo( Program, CL_PROGRAM_BINARIES, size, &binary, NULL );
PrintCLError( status, "clGetProgramInfo (2):" );

FILE * fpbin = fopen( CL_BINARY_NAME, "wb" );
if( fpbin == NULL )
{
    fprintf( stderr, "Cannot create \"%s\\n\", CL_BINARY_NAME );
}
else
{
    fwrite( binary, 1, size, fpbin );
    fclose( fpbin );
}
delete [ ] binary;
```

This binary can then be used in a call to
clCreateProgramWithBinary()

mjb - March 16, 2015

NVIDIA OpenCL Assembly Language Sample

```
ld.global.v4.f32    {%f188, %f189, %f190, %f191}, [%r1];    // load dPobj[ gid ]
ld.global.v4.f32    {%f156, %f157, %f158, %f159}, [%r2];    // load dVel[ gid ]

mov.f32            %f17, 0f3DCCCCCD;                        // DT

fma.rn.f32         %f248, %f156, %f17, %f188;              // (p + v*DT).x
fma.rn.f32         %f249, %f157, %f17, %f189;              // (p + v*DT).y
fma.rn.f32         %f250, %f158, %f17, %f190;              // (p + v*DT).z

mov.f32            %f18, 0fBD48B43B;                        // .5 * G.y * DT * DT
mov.f32            %f19, 0f00000000;                        // 0., for .x and .z

add.f32            %f256, %f248, %f19;                      // (p + v*DT).x + 0.
add.f32            %f257, %f249, %f18;                      // (p + v*DT).y + .5 * G.y * DT * DT
add.f32            %f258, %f250, %f19;                      // (p + v*DT).z + 0.

mov.f32            %f20, 0fBF7AE148;                        // G.y * DT

add.f32            %f264, %f156, %f19;                      // v.x + 0.
add.f32            %f265, %f157, %f20;                      // v.y + G.y * DT
add.f32            %f266, %f158, %f19;                      // v.z + 0.
```

mjb - March 16, 2015

Things Learned from Examining OpenCL Assembly Language on an NVIDIA GTX 480

- The points, vectors, and colors were typedef'ed as float4's, but the compiler realized that they were being used as float3's, and didn't bother with the 4th element.
- The floatn's were not SIMD'ed. (We actually knew this already, since NVIDIA doesn't supported vector operations in their GPUs.) There is still an advantage in coding this way, even if just for readability.
- The function calls were all in-lined. (This makes sense – the OpenCL spec says “no recursion”, which implies “no stack”, which would make function calls difficult.)
- Defining G, DT, and Sphere1 as **constant** memory types was a mistake. It got the correct results, but the compiler didn't take advantage of them being constants. Changing them to type **const** threw compiler errors because of their global scope. Changing them to **const** and moving them into the body of the kernel function Particle *did* result in compiler optimizations.
- $\sqrt{x^2+y^2+z^2}$ is amazingly involved. I can only hope that there is a good reason. Use `fast_sqrt()`, `fast_normalize()`, and `fast_length()` when you can.
- The compiler did not do a good job with expressions-in-common. I had really hoped it would figure out that detecting if a point was in a sphere and determining the unitized surface normal at that point were mostly the same operation, but it didn't.
- There is a 4-argument floating-point-multiply-and-add instruction ($d = a + b*c$), which the compiler took great advantage of.

mjb - March 16, 2015