


## Performing Reductions in OpenCL

Mike Bailey

mjb@cs.oregonstate.edu

Oregon State University



Oregon State University  
Computer Graphics

opengl.reduction.pptx


mjb - February 18, 2015

## Here's the Problem We are Trying to Solve

Like the first demo program, we are piecewise multiplying two arrays. Unlike the first demo program, we want to then add up all the products and return the sum.

$$\begin{aligned} A * B &\rightarrow \text{prods} \\ \sum \text{prods} &\rightarrow C \end{aligned}$$

numItems = 8;

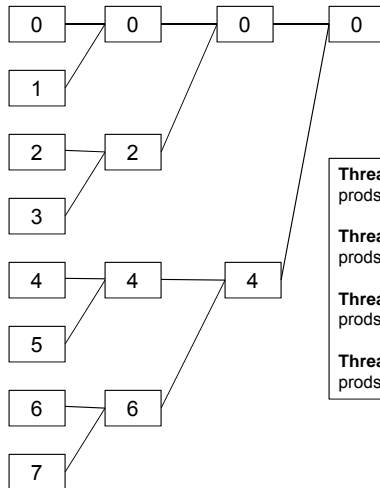


Oregon State University  
Computer Graphics

mjb - February 18, 2015

## Reduction Takes Place in a Single Work-Group

numItems = 8;



If we had 8 work-items in a work-group, we would like the threads in each work-group to execute the following instructions . . .

<b>Thread #0:</b> prods[ 0 ] += prods[ 1 ] ;	<b>Thread #0:</b> prods[ 0 ] += prods[ 2 ] ;	<b>Thread #0:</b> prods[ 0 ] += prods[ 4 ] ;
<b>Thread #2:</b> prods[ 2 ] += prods[ 3 ] ;		
<b>Thread #4:</b> prods[ 4 ] += prods[ 5 ] ;	<b>Thread #4:</b> prods[ 4 ] += prods[ 6 ] ;	
<b>Thread #6:</b> prods[ 6 ] += prods[ 7 ] ;		

. . . but in a more general way than writing them all out by hand.



Oregon State University  
Computer Graphics

mjb - February 18, 2015

## Here's What You Would Change in your Host Program

```
size_t numWorkGroups = NUM_ELEMENTS / LOCAL_SIZE;

. . .

float * hA = new float [ NUM_ELEMENTS ];
float * hB = new float [ NUM_ELEMENTS ];
float * hC = new float [ numWorkGroups ];
size_t abSize = NUM_ELEMENTS * sizeof(float);
size_t cSize = numWorkGroups * sizeof(float);

. . .

cl_mem dA = clCreateBuffer( context, CL_MEM_READ_ONLY, abSize, NULL, &status );
cl_mem dB = clCreateBuffer( context, CL_MEM_READ_ONLY, abSize, NULL, &status );
cl_mem dC = clCreateBuffer( context, CL_MEM_WRITE_ONLY, cSize, NULL, &status );

. . .

status = clEnqueueWriteBuffer( cmdQueue, dA, CL_FALSE, 0, abSize, hA, 0, NULL, NULL );
status = clEnqueueWriteBuffer( cmdQueue, dB, CL_FALSE, 0, abSize, hB, 0, NULL, NULL );

. . .

cl_kernel kernel = clCreateKernel( program, "ArrayMultReduce", &status );

. . .

status = clSetKernelArg( kernel, 0, sizeof(cl_mem), &dA );
status = clSetKernelArg( kernel, 1, sizeof(cl_mem), &dB );
status = clSetKernelArg( kernel, 2, sizeof(float), NULL ); // local "prods" array - one per work-item
status = clSetKernelArg( kernel, 3, sizeof(cl_mem), &dC );
```

$A * B \rightarrow \text{prods}$   
 $\sum \text{prods} \rightarrow C$



mjb - February 18, 2015

## The Arguments to the Kernel

```
status = clSetKernelArg( kernel, 0, sizeof(cl_mem), &dA );
status = clSetKernelArg( kernel, 1, sizeof(cl_mem), &dB );
status = clSetKernelArg( kernel, 2, sizeof(float), NULL ); // local "prods" array – one per work-item
status = clSetKernelArg( kernel, 3, sizeof(cl_mem), &dC );
```

```
kernel void
ArrayMultReduce( global const float *dA, global const float *dB, local float *prods, global float *dC )
{
    int gid      = get_global_id( 0 );
    int tnum     = get_local_id( 0 ); // thread number
    int wgNum    = get_group_id( 0 ); // work-group number
    int numItems = get_local_size( 0 );

    prods[ tnum ] = dA[ gid ] * dB[ gid ]; // multiply the two arrays together

    // now add them up – come up with one sum per work-group
    // it is a big performance benefit to do it here while "prods" is still available – and is local
    // it would be a performance hit to pass "prods" back to the host then bring it back to the device for reduction
}
```

$A * B \rightarrow \text{prods}$



Oregon State University  
Computer Graphics

mjb – February 18, 2015

## Reduction Takes Place in a Single Work-Group Each work-item is run by a single thread

<b>Thread #0:</b> <code>prods[ 0 ] += prods[ 1 ];</code>	<b>Thread #0:</b> <code>prods[ 0 ] += prods[ 2 ];</code>	<b>Thread #0:</b> <code>prods[ 0 ] += prods[ 4 ];</code>
<b>Thread #2:</b> <code>prods[ 2 ] += prods[ 3 ];</code>		<code>offset = 4</code> <code>mask = 7;</code>
<b>Thread #4:</b> <code>prods[ 4 ] += prods[ 5 ];</code>	<b>Thread #4:</b> <code>prods[ 4 ] += prods[ 6 ];</code>	
<b>Thread #6:</b> <code>prods[ 6 ] += prods[ 7 ];</code>	<code>offset = 2</code> <code>mask = 3;</code>	

`offset = 1`  
`mask = 1;`

`numItems = 8;`



Oregon State University  
Computer Graphics

mjb – February 18, 2015

## Reduction Takes Place in a Single Work-Group

Each work-item is run by a single thread

**Thread #0:**  
prods[ 0 ] += prods[ 1 ];

**Thread #2:**  
prods[ 2 ] += prods[ 3 ];

**Thread #4:**  
prods[ 4 ] += prods[ 5 ];

**Thread #6:**  
prods[ 6 ] += prods[ 7 ];

offset = 1;  
mask = 1;

numItems = 8;

**Thread #0:**  
prods[ 0 ] += prods[ 2 ];

**Thread #4:**  
prods[ 4 ] += prods[ 6 ];

offset = 2;  
mask = 3;

**Thread #0:**  
prods[ 0 ] += prods[ 4 ];

offset = 4;  
mask = 7;

```
kernel void
ArrayMultReduce( ... )
{
    int gid    = get_global_id( 0 );
    int tnum   = get_local_id( 0 ); // thread number
    int wgNum  = get_group_id( 0 ); // work-group number
    int numItems = get_local_size( 0 );

    prods[ tnum ] = dA[ gid ] * dB[ gid ];
}
```

```
// all threads execute this code simultaneously:
for( int offset = 1; offset < numItems; offset *= 2 )
{
    int mask = 2*offset - 1;
    barrier( CLK_LOCAL_MEM_FENCE );
    if( ( tnum & mask ) == 0 )
    {
        prods[ tnum ] += prods[ tnum + offset ];
    }
}
```

```
barrier( CLK_LOCAL_MEM_FENCE );
if( tnum == 0 )
    dC[ wgNum ] = prods[ 0 ];
```



Oregon State University  
Computer Graphics

$\sum \text{prods} \rightarrow C$

## And, Finally, in your Host Program

```
status = clEnqueueReadBuffer( cmdQueue, dC, CL_TRUE, 0, cSize, hC, 0, NULL, NULL );

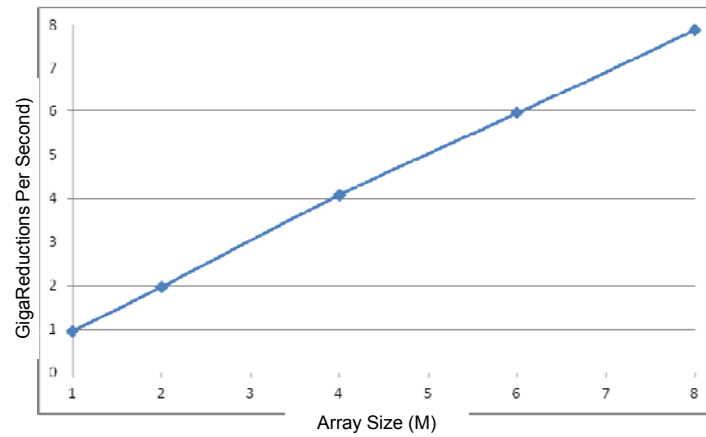
float sum = 0.;
for( int i = 0; i < numWorkgroups; i++ )
{
    sum += hC[ i ];
}
```



Oregon State University  
Computer Graphics

mjb - February 18, 2015

### Reduction Performance: Work-Group Size = 32



Oregon State University  
Computer Graphics

mjb - February 18, 2015