

## OpenMP Case Study: Trapezoid Integration Example

Mike Bailey  
mjb@cs.oregonstate.edu

Oregon State University




Oregon State University  
Computer Graphics

trapezoid.pptx

mjb - March 12, 2015

Find the area under the curve  $y = \sin(x)$   
for  $0 \leq x \leq \pi$   
using the Trapezoid Rule



Exact answer:  $\int_0^{\pi} (\sin x) dx = -\cos x \Big|_0^{\pi} = 2.0$



Oregon State University  
Computer Graphics

mjb - March 12, 2015

## Don't do it this way !

```
const double A = 0.;
const double B = M_PI;

double dx = ( B - A ) / (float) ( numSubdivisions - 1 );
double sum = ( Function( A ) + Function( B ) ) / 2.;

omp_set_num_threads( numThreads );

#pragma omp parallel for default(none),shared(dx,sum)
for( int i = 1; i < numSubdivisions - 1; i++ )
{
    double x = A + dx * (float) i;
    double f = Function( x );
    sum += f;
}

sum *= dx;
```

### Assembly code:

```
Load sum
Add f
Store sum
```


What if the scheduler  
decides to switch  
threads right here?



Oregon State University  
Computer Graphics

mjb - March 12, 2015

The answer should be 2.0, but in 30 trials, it's not even close.  
And, the answers aren't even consistent. Why?



Oregon State University  
Computer Graphics

mjb - March 12, 2015

### Do it this way !

```
const double A = 0.;
const double B = M_PI;

double dx = ( B - A ) / (float) ( numSubdivisions - 1 );

omp_set_num_threads( numThreads );

double sum = ( Function( A ) + Function( B ) ) / 2.;

#pragma omp parallel for default(none),shared(dx) reduction(+:sum)
for( int i = 1; i < numSubdivisions - 1; i++ )
{
    double x = A + dx * (float) i;
    double f = Function( x );
    sum += f;
}


sum *= dx;
```



Oregon State University  
Computer Graphics

mjb - March 12, 2015


### MegaFunctionEvaluations Per Second vs. Number of Subdivisions



Oregon State University  
Computer Graphics

mjb - March 12, 2015

## MegaFunctionEvaluations Per Second vs. Number of Threads



OSU

Oregon State University  
Computer Graphics

mjb - March 12, 2015

## Speed of using Reduction vs. Atomic vs. Critical

```
#pragma omp parallel for shared(dx) reduction(+:sum)
for( int i = 0; i < numSubdivisions; i++ )
{
    double x = A + dx * (float) i;
    double f = Function( x );
    sum += f;
}
```

1

```
#pragma omp parallel for shared(dx)
for( int i = 0; i < numSubdivisions; i++ )
{
    double x = A + dx * (float) i;
    double f = Function( x );
    #pragma omp atomic
    sum += f;
}
```

2

```
#pragma omp parallel for shared(dx)
for( int i = 0; i < numSubdivisions; i++ )
{
    double x = A + dx * (float) i;
    double f = Function( x );
    #pragma omp critical
    sum += f;
}
```

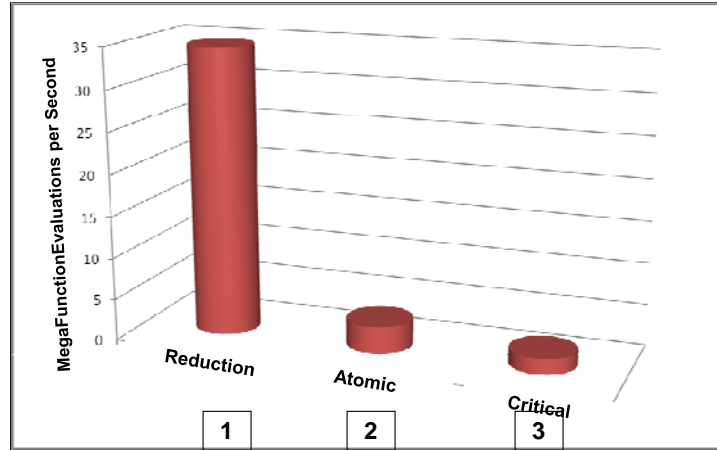
3

OSU

Computer Graphics

mjb - March 12, 2015



## Speed of using Reduction vs. Atomic vs. Critical



Oregon State University  
Computer Graphics

mjb - March 12, 2015

## Two Reasons Why Reduction is so Much Better in this Case



1. Reduction automatically creates a temporary private variable for each thread's running sum. Each thread adding into its running sum doesn't interfere with any other thread adding into its running sum, and so threads don't need to slow down to get out of the way of each other.
2. Reduction automatically creates a binary tree structure, like this, to add the  $N$  running sums in  $\log_2 N$  time instead of  $N$  time.



Oregon State University  
Computer Graphics

mjb - March 12, 2015