

Parallel Programming using OpenMP

Mike Bailey

mjb@cs.oregonstate.edu

Oregon State University



Oregon State University
Computer Graphics

openmp.pptx

mjb - March 12, 2015

OpenMP Multithreaded Programming

- OpenMP stands for "Open Multi-Processing"
- OpenMP is a multi-vendor* standard to perform shared-memory multithreading
- OpenMP uses the fork-join model
- OpenMP is both directive- and library-based
- OpenMP threads share a single executable, global memory, and heap (malloc, new)
- Each OpenMP thread has its own stack (function arguments, local variables)
- Using OpenMP requires no dramatic code changes
- OpenMP probably gives you the biggest multithread benefit per amount of work you have to put in to using it

**Much of your use of OpenMP will be accomplished by issuing C/C++
"pragmas" to tell the compiler how to build the threads into the executable**

```
#pragma omp directive [clause]
```



Oregon State Un
Computer Gra

* AMD, Fujitsu, HP, IBM, Intel, Microsoft, NEC, NVIDIA, Oracle, Texas Instruments, VMWare, ...

mjb - March 12, 2015

What OpenMP Isn't:

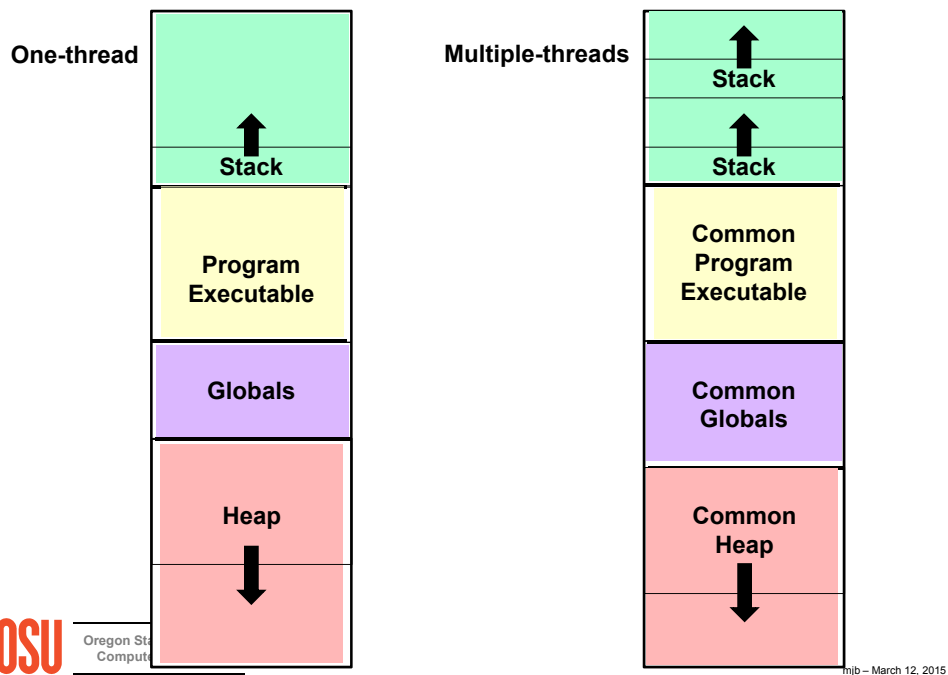
- OpenMP doesn't check for data dependencies, data conflicts, deadlocks, or race conditions. You are responsible for avoiding those yourself
- OpenMP doesn't check for non-conforming code sequences
- OpenMP doesn't guarantee *identical* behavior across vendors or hardware
- OpenMP doesn't guarantee the *order* in which threads execute, just that they do execute
- OpenMP is not overhead-free
- OpenMP does not prevent you from writing false-sharing code (in fact, it makes it really easy)



Oregon State University
Computer Graphics

mjb - March 12, 2015

Memory Allocation in a Multithreaded Program



Oregon State University
Computer Graphics

mjb - March 12, 2015

Using OpenMP in Linux

```
g++ -o proj proj.cpp -O3 -lm -fopenmp
```

```
icpc -o proj proj.cpp -O3 -lm -openmp -align -qopt-report=3 -qopt-report-phase=vec
```

Using OpenMP in Microsoft Visual Studio

1. Go to the Project menu → Project Properties
2. Change the setting Configuration Properties → C/C++ → Language → OpenMP Support to **"Yes (/openmp)"**

Seeing if OpenMP is Supported on Your System

```
#ifndef _OPENMP
fprintf( stderr, "OpenMP is not supported – sorry!\n" );
#endif
```



Oregon State University
Computer Graphics

mjb – March 12, 2015

Number of OpenMP threads

Two ways to specify how many OpenMP threads you want to have available:

1. Set the OMP_NUM_THREADS environment variable
2. Call `omp_set_num_threads(num);`

Asking how many cores this program has access to:

```
num = omp_get_num_procs( );
```

Setting the number of threads to the exact number of cores available:

```
num = omp_set_num_threads( omp_get_num_procs( ) );
```

Asking how many OpenMP threads this program is using right now:

```
num = omp_get_num_threads( );
```

Asking which thread this one is:

```
me = omp_get_thread_num( );
```



Oregon State University
Computer Graphics

mjb – March 12, 2015

Creating an OpenMP Team of Threads

```
#pragma omp parallel default(none)
{
    ...
}
```

This creates a team of threads



Each thread would then execute all lines of code in this block.

Try this, just for fun:

```
omp_set_num_threads( 4 );
#pragma omp parallel default(none)
printf( "Hello, World, from thread #%%d ! \n" , omp_get_thread_num( ) );
```

Hint: run it several times in a row. What do you see? Why?



Oregon State University
Computer Graphics

mjb - March 12, 2015

Uh-oh...

```
Hello, World, from thread #6 !
Hello, World, from thread #1 !
Hello, World, from thread #7 !
Hello, World, from thread #5 !
Hello, World, from thread #4 !
Hello, World, from thread #3 !
Hello, World, from thread #2 !
Hello, World, from thread #0 !
```

```
Hello, World, from thread #0 !
Hello, World, from thread #7 !
Hello, World, from thread #4 !
Hello, World, from thread #6 !
Hello, World, from thread #1 !
Hello, World, from thread #3 !
Hello, World, from thread #5 !
Hello, World, from thread #2 !
```

```
Hello, World, from thread #2 !
Hello, World, from thread #5 !
Hello, World, from thread #0 !
Hello, World, from thread #7 !
Hello, World, from thread #1 !
Hello, World, from thread #3 !
Hello, World, from thread #4 !
Hello, World, from thread #6 !
```

```
Hello, World, from thread #1 !
Hello, World, from thread #3 !
Hello, World, from thread #5 !
Hello, World, from thread #2 !
Hello, World, from thread #4 !
Hello, World, from thread #7 !
Hello, World, from thread #6 !
Hello, World, from thread #0 !
```



Oregon State University
Computer Graphics




mjb - March 12, 2015

Creating OpenMP threads in Loops

```
#include <omp.h>
int i;
#pragma omp parallel for default(none),private(i)
for( i = 0; i < num; i++ )
{
    ...
}
```

This creates a
team of threads



There is an “**implied
barrier**” at the end

This tells the compiler to parallelize the for loop into multiple threads, and to give each thread its own personal copy of the variable *i*. But, you don’t have to do this for variables defined in the loop body:

```
#pragma omp parallel for default(none)
for( int i = 0; i < num; i++ )
{
    ...
}
```



Oregon State University
Computer Graphics

mjb – March 12, 2015

OpenMP for-Loop Rules

```
#pragma omp parallel for default(none), shared(...), private(...)
for( int index = start ; index terminate condition; index changed )
```

- The *index* must be an *int* or a *pointer*
- The *start* and *end* conditions must have compatible types
- Neither the *start* nor the *end* conditions can be changed during the execution of the loop
- The *index* can only be modified by the “changed” expression (i.e., not modified inside the loop itself)
- There can be no between-loop data dependencies

This is the probably the biggest parallel benefit per programming effort !



Oregon State University
Computer Graphics

mjb – March 12, 2015

OpenMP For-Loop Rules

```
for( index = start ;      index <  end      index <= end      index >  end      index >= end      index++      ++index      index--      --index      index += incr      index = index + incr      index = incr + index      index -= decr      index = index - decr      )
```



Oregon State University
Computer Graphics

mjb - March 12, 2015

OpenMP Directive Data Types

I recommend that you use:

default(none)

in all your OpenMP directives. This will force you to explicitly flag all of your inside variables as shared or private. This will help prevent mistakes.

private(x)

Means that each thread will have its own copy of the variable x

shared(x)

Means that each thread will share a common x. This is potentially dangerous.

Example:

```
#pragma omp parallel for default(none),private(i,j),shared(x)
```



Oregon State University
Computer Graphics



mjb - March 12, 2015

OpenMP Allocation of Work to Threads

Static Threads

- All work is allocated and assigned at runtime

Dynamic Threads

- Consists of one Master and a pool of threads
- The pool is assigned some of the work at runtime, but not all of it
- When a thread from the pool becomes idle, the Master gives it a new assignment
- "Round-robin assignments"

OpenMP Scheduling

`schedule(static [,chunksize])`

`schedule(dynamic [,chunksize])`

Defaults to static

chunksize defaults to 1

In static, the iterations are assigned to threads before the loop starts



Oregon State University
Computer Graphics

mjb - March 12, 2015

OpenMP Allocation of Work to Threads

```
#pragma omp parallel for default(none),schedule(static,chunksize)
for( int index = 0 ; index < 12 ; index++ )
```

Static,1

0	0,3,6,9
1	1,4,7,10
2	2,5,8,11

chunksize = 1

Each thread is assigned one iteration, then the assignments start over

Static,2

0	0,1,6,7
1	2,3,8,9
2	4,5,10,11

chunksize = 2

Each thread is assigned two iterations, then the assignments start over

Static,4

0	0,1,2,3
1	4,5,6,7
2	8,9,10,11

chunksize = 4

Each thread is assigned four iterations, then the assignments start over



Oregon State University
Computer Graphics



mjb - March 12, 2015

Arithmetic Operations Among Threads – A Problem

```
#pragma omp parallel for private(myPartialSum),shared(sum)
for( int i = 0; i < N; i++ )
{
    float myPartialSum = ...

    sum = sum + myPartialSum;
}
```

- There is no guarantee when each thread will execute this line correctly
- There is not even a guarantee that each thread will finish this line before some other thread interrupts it
- This is non-deterministic !

Assembly code:

```
Load sum
Add myPartialSum
Store sum
```

What if the scheduler decides to switch threads right here?



Oregon State University
Computer Graphics

Conclusion: Don't do it this way!

mjb - March 12, 2015

Here's a trapezoid integration example (covered in another section).


The partial sums are added up, as shown on the previous page.

The integration was done 30 times.

The answer is supposed to be exactly 2.

None of the 30 answers is even close.

And, not only are the answers bad, they are not even consistently bad!



Oregon State University
Computer Graphics

Don't do it this way!

mjb - March 12, 2015

Arithmetic Operations Among Threads – Three Solutions

#pragma omp atomic

sum = sum + myPartialSum;

1

- Fixes the non-deterministic problem
- But, serializes the code
- Operators include +, -, *, /, ++, --, >>, <<, ^, |
- Operators include +=, -=, *=, /=, etc.

#pragma omp critical

sum = sum + myPartialSum;

2

- Also fixes it
- But, serializes the code


#pragma omp parallel for reduction(+:sum),private(myPartialSum)

...

sum = sum + myPartialSum;

3

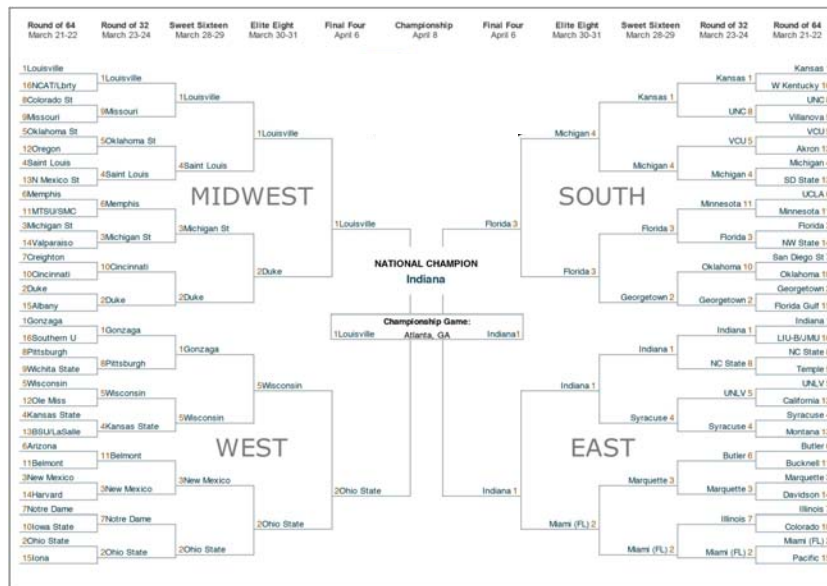
- Performs (sum,product,and,or,...) in $O(\log_2 N)$ time instead of $O(N)$
- Operators include +, -, *, /, ++, --
- Operators include +=, -=, *=, /=
- Operators include ^=, |=, &=



Oregon State University
Computer Graphics

mjb – March 12, 2015

If You Understand Basketball Brackets, You'll Understand Reduction




Oregon State University
Computer Graphics

Source: ESPN

mjb – March 12, 2015

Reduction vs. Atomic vs. Critical



Oregon State University
Computer Graphics



mjb - March 12, 2015

Synchronization

Mutual Exclusion Locks (Mutexes)

```
omp_init_lock( omp_lock_t * );
omp_set_lock( omp_lock_t * ); ← Blocks if the lock is not available
                                Then sets it and returns when it is available
omp_unset_lock( omp_lock_t * );
omp_test_lock( omp_lock_t * ); ← If the lock is not available, returns 0
                                If the lock is available, sets it and returns !0
```

(omp_lock_t is really an array of 4 unsigned chars)

Critical sections

```
#pragma omp critical
    Restricts execution to one thread at a time
```

```
#pragma omp single
    Restricts execution to a single thread ever
```

Barriers

```
#pragma omp barrier
    Forces all threads to wait here until all threads arrive
```

(Note: there is an implied barrier after parallel for loops and OpenMP sections, unless the *nowait* clause is used)



Oregon State University
Computer Graphics

mjb - March 12, 2015

Synchronization Examples

```
omp_lock_t      Sync;
...
omp_init_lock( &Sync );

...

omp_set_lock( &Sync );
    << code that needs the mutual exclusion >>
omp_unset_lock( &Sync );

...

while( omp_test_lock( &Sync ) == 0 )
{
    DoSomeUsefulWork( );
}
```



Oregon State University
Computer Graphics



mjb - March 12, 2015

Creating Sections of OpenMP Code

Sections are consecutive, independent blocks of code

```
#pragma omp parallel sections
{
    #pragma omp section
    {
        ...
    }

    #pragma omp section
    {
        ...
    }
}
```

Each section is executed by *some* thread (not necessarily its *own* thread)

There is an implied barrier at the end



Oregon State University
Computer Graphics

mjb - March 12, 2015

OpenMP Tasks

- An OpenMP task is a single line of code or a structured block which is immediately assigned to *one thread in the current thread team*
- If the *ties* clause is used, it is assigned to the same thread. If the *unties* clause is used, it can be assigned to any thread.
- The task can be executed immediately, or it can be placed on its thread's list of things to do.
- If the *if* clause is used and the argument evaluates to 0, then the task is executed immediately, superceding whatever else that thread is doing.
- There has to be an existing parallel thread team for this to be effective. Otherwise one thread ends up doing all tasks.
- One of the best uses of this is to make a function call. That function then runs concurrently until it completes.

```
#pragma omp task
Watch_For_Internet_Input( );
```

You can create a task barrier with:

```
#pragma omp taskwait
```

These are very much like OpenMP **Sections**, but Sections are more static, that is, they are setup when you write the code, whereas **Tasks** can be created anytime, and in any number, under control of your program's logic.




Oregon State University
Computer Graphics

mjb - March 12, 2015

OpenMP Task Example: Processing each element of a linked list

```
#pragma omp parallel
{
    #pragma omp single default(none)
    {
        element *p = listHead;
        while( p != NULL )
        {
            #pragma omp task
            Process( p );

            p = p->next;
        }
    }
}
```




Oregon State University
Computer Graphics

mjb - March 12, 2015

Single Program Multiple Data (SPMD) in OpenMP

```
#define NUM      1000000
float A[NUM], B[NUM], C[NUM];

total = omp_get_num_threads( );
#pragma omp parallel default(none),private(me),shared(total)
{
    me = omp_get_thread_num( );
    DoWork( me, total );
}
```



```
void DoWork( int me, int total )
{
    int first = NUM * me / total;
    int last = NUM * (me+1)/total - 1;
    for( int i = first; i <= last; i++ )
    {
        C[i] = A[i] * B[i];
    }
}
```



Oregon State University
Computer Graphics

mjb - March 12, 2015