



The Intel Xeon Phi

Mike Bailey

mjb@cs.oregonstate.edu

Oregon State University

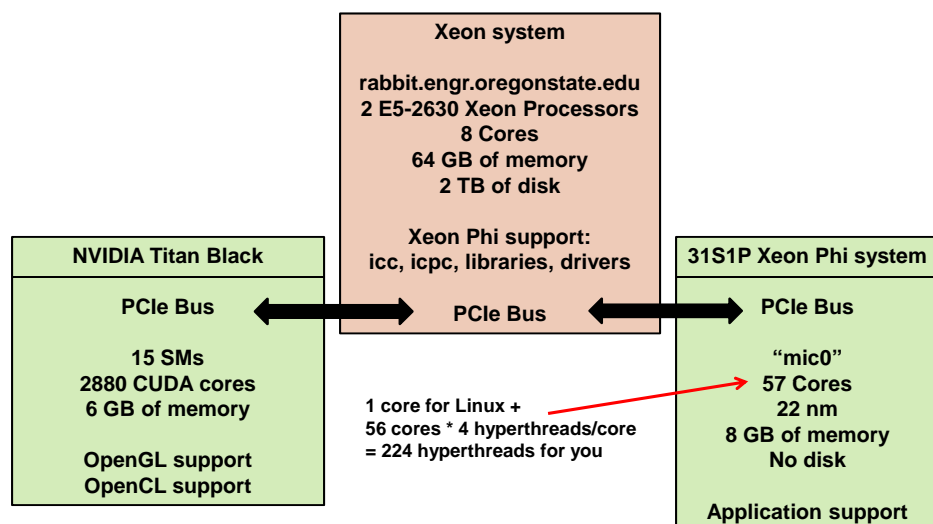


Oregon State University
Computer Graphics

OSU_Xeon_Phi.pptx

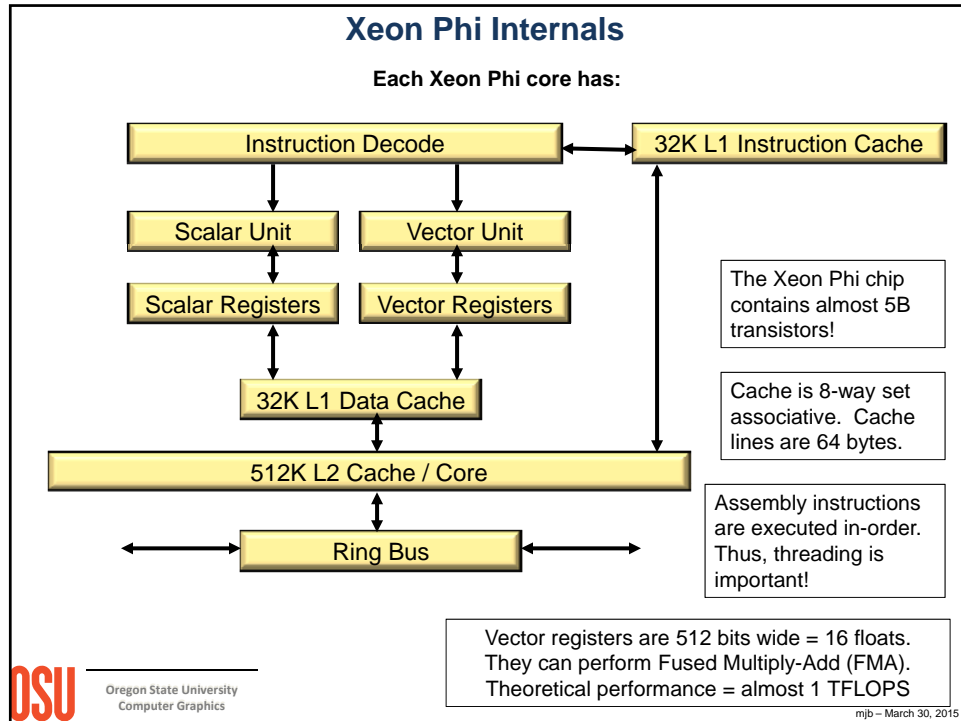
mjb - March 30, 2015

Setup



Oregon State University
Computer Graphics

mjb - March 30, 2015



Xeon Phi Peak Performance

Clock freq x # cores x # vector lanes x 2 FMA / 2 cycles to decode =

1.091 GHz x 56 x 16 x 2 / 2 =

0.98 TFLOPS

FMA stands for "Fused Multiply+Add".

It allows the operation:

$$d = a * b + c;$$

to be performed in the same amount of time as:

$$d = a * b;$$

OSU Oregon State University Computer Graphics

mjb - March 30, 2015

Getting to *rabbit* and setting up your account

Lowercase letter 'L'

To login to *rabbit*:

```
ssh rabbit.engr.oregonstate.edu -l yourengusername
```

Put this in your *rabbit* account's `.cshrc` :

```
setenv INTEL_LICENSE_FILE 28518@lic.engr.oregonstate.edu
setenv SINK_LD_LIBRARY_PATH /nfs/guille/a2/rh80apps/intel/studio.2013-sp1/composer_xe_2015.0.090/compiler/lib/mic/
setenv ICCPATH /nfs/guille/a2/rh80apps/intel/studio.2013-sp1/composer_xe_2015/bin/
set path=( $path $ICCPATH )
source /nfs/guille/a2/rh80apps/intel/studio.2013-sp1/bin/iccvars.csh intel64
```

Then activate these values like this:

```
source .cshrc
```

(These will be activated automatically the next time you login.)

To verify that the Xeon Phi card is there:

```
ping mic0
```

To see the Xeon Phi card characteristics:

```
micinfo
```

To run some operational tests on the Xeon Phi:

```
miccheck
```



mjb - March 30, 2015

Running *ping*

rabbit 150% ping mic0

```
PING rabbit-mic0.engr.oregonstate.edu (172.31.1.1) 56(84) bytes of data:
64 bytes from rabbit-mic0.engr.oregonstate.edu (172.31.1.1): icmp_seq=1 ttl=64 time=290 ms
64 bytes from rabbit-mic0.engr.oregonstate.edu (172.31.1.1): icmp_seq=2 ttl=64 time=0.385 ms
64 bytes from rabbit-mic0.engr.oregonstate.edu (172.31.1.1): icmp_seq=3 ttl=64 time=0.242 ms
64 bytes from rabbit-mic0.engr.oregonstate.edu (172.31.1.1): icmp_seq=4 ttl=64 time=0.230 ms
64 bytes from rabbit-mic0.engr.oregonstate.edu (172.31.1.1): icmp_seq=5 ttl=64 time=0.225 ms
64 bytes from rabbit-mic0.engr.oregonstate.edu (172.31.1.1): icmp_seq=6 ttl=64 time=0.261 ms
```



Oregon State University
Computer Graphics

mjb - March 30, 2015

Running *micinfo*

rabbit 151% micinfo

MicInfo Utility Log
Created Mon Jan 12 10:21:07 2015

System Info

HOST OS : Linux
OS Version : 2.6.32-504.3.3.el6.x86_64
Driver Version : 3.4.2-1
MPSS Version : 3.4.2
Host Physical Memory : 65859 MB

Device No: 0, Device Name: mic0

Version

Flash Version : 2.1.02.0390
SMC Firmware Version : 1.16.5078
SMC Boot Loader Version : 1.8.4326
uOS Version : 2.6.38.8+mpss3.4.2
Device Serial Number : ADKC31600731

Board

Vendor ID : 0x8086
Device ID : 0x225e
Subsystem ID : 0x2500
Coprocessor Stepping ID : 3
PCIe Width : Insufficient Privileges
PCIe Speed : Insufficient Privileges
PCIe Max payload size : Insufficient Privileges
PCIe Max read req size : Insufficient Privileges

Coprocessor Model : 0x01
Coprocessor Model Ext : 0x00
Coprocessor Type : 0x00
Coprocessor Family : 0x0b
Coprocessor Family Ext : 0x00
Coprocessor Stepping : B1
Board SKU : B1PRQ-31S1P
ECC Mode : Enabled
SMC HW Revision : Product 300W Passive CS

Cores

Total No of Active Cores : 57
Voltage : 1089000 uV
Frequency : 1100000 kHz

Thermal

Fan Speed Control : N/A
Fan RPM : N/A
Fan PWM : N/A
Die Temp : 40 C

GDDR

GDDR Vendor : Elpida
GDDR Version : 0x1
GDDR Density : 2048 Mb
GDDR Size : 7936 MB
GDDR Technology : GDDR5
GDDR Speed : 5.000000 GT/s
GDDR Frequency : 2500000 kHz
GDDR Voltage : 1501000 uV



Oregon State University
Computer Graphics

mjb - March 30, 2015

Running *miccheck*

rabbit 152% miccheck

MicCheck 3.4.2-r1
Copyright 2013 Intel Corporation All Rights Reserved

Executing default tests for host

Test 0: Check number of devices the OS sees in the system ... pass
Test 1: Check mic driver is loaded ... pass
Test 2: Check number of devices driver sees in the system ... pass
Test 3: Check mpssd daemon is running ... Pass

Executing default tests for device: 0

Test 4 (mic0): Check device is in online state and its postcode is FF ... pass
Test 5 (mic0): Check ras daemon is available in device ... pass
Test 6 (mic0): Check running flash version is correct ... pass
Test 7 (mic0): Check running SMC firmware version is correct ... pass

Status: OK



Oregon State University
Computer Graphics

mjb - March 30, 2015

Running *micsmc, I*

rabbit 153% micsmc -a

mic0 (info):

Device Series: Intel(R) Xeon Phi(TM) coprocessor x100 family
 Device ID: 0x225e
 Number of Cores: 57
 OS Version: 2.6.38.8+mpss3.4.2
 Flash Version: 2.1.02.0390
 Driver Version: 3.4.2-1 (root@rabbit.engr.oregonstate.edu)
 Stepping: 0x3
 Substepping: 0x0

mic0 (temp):

Cpu Temp: 44.00 C
 Memory Temp: 28.00 C
 Fan-In Temp: 24.00 C
 Fan-Out Temp: 28.00 C
 Core Rail Temp: 29.00 C
 Uncore Rail Temp: 29.00 C
 Memory Rail Temp: 29.00 C

mic0 (freq):

Core Frequency: 1.10 GHz
 Total Power: 92.00 Watts
 Low Power Limit: 283.00 Watts
 High Power Limit: 337.00 Watts
 Physical Power Limit: 357.00 Watts

mic0 (mem):

Free Memory: 7347.64 MB
 Total Memory: 7698.83 MB
 Memory Usage: 351.18 MB



Oregon State University
Computer Graphics

mjb - March 30, 2015

Running *micsmc, II*

mic0 (cores):

Device Utilization: User: 0.00%, System: 0.09%, Idle: 99.91%
 Per Core Utilization (57 cores in use)

Core #1: User: 0.00%, System: 0.27%, Idle: 99.73%
 Core #2: User: 0.00%, System: 0.27%, Idle: 99.73%
 Core #3: User: 0.00%, System: 0.00%, Idle: 100.00%
 Core #4: User: 0.00%, System: 0.00%, Idle: 100.00%
 Core #5: User: 0.00%, System: 0.00%, Idle: 100.00%
 Core #6: User: 0.00%, System: 0.00%, Idle: 100.00%
 Core #7: User: 0.00%, System: 0.00%, Idle: 100.00%
 Core #8: User: 0.00%, System: 0.27%, Idle: 99.73%
 Core #9: User: 0.00%, System: 0.00%, Idle: 100.00%
 Core #10: User: 0.00%, System: 0.27%, Idle: 99.73%

...

Core #50: User: 0.00%, System: 0.00%, Idle: 100.00%
 Core #52: User: 0.00%, System: 0.27%, Idle: 99.73%
 Core #53: User: 0.00%, System: 0.00%, Idle: 100.00%
 Core #54: User: 0.00%, System: 0.27%, Idle: 99.73%
 Core #55: User: 0.00%, System: 0.00%, Idle: 100.00%
 Core #56: User: 0.00%, System: 0.27%, Idle: 99.73%
 Core #57: User: 0.00%, System: 0.54%, Idle: 99.46%



Oregon State University
Computer Graphics



mjb - March 30, 2015

Cross-compiling and running from *rabbit*

To compile on *rabbit* for *rabbit*:

```
icpc -o try try.cpp -O3 -lm -openmp -align -qopt-report=3 -qopt-report-phase=vec
```

To cross-compile on *rabbit* for the Xeon Phi:

```
icpc -mmic -o try try.cpp -O3 -lm -openmp -align -qopt-report=3 -qopt-report-phase=vec
```

Note: the summary of vectorization success or failure is in a *.optvec file

To execute on the Xeon Phi, type this on *rabbit*:

```
micnativeloadex try
```



Oregon State University
Computer Graphics

mjb - March 30, 2015

Gaining Access to the Cores, I

```
#pragma omp parallel for
for( int i = 0; i < N; i++ )
    C[i] = A[i] * B[i];
```

```
float sum = 0.;
#pragma omp parallel for reduction(+:sum)
for( int i = 0; i < N; i++ )
    sum += A[i] * B[i];
```

```
icpc -mmic -o try try.cpp -O3 -lm -openmp -align -qopt-report=3 -qopt-report-phase=vec
micnativeloadex try
```



Oregon State University
Computer Graphics

mjb - March 30, 2015

Gaining Access to the Cores, II

```
#pragma omp parallel sections
#pragma omp section
...
#pragma omp section
...
```

```
#pragma omp task
...
```

```
icpc -mmic -o try try.cpp -O3 -m -openmp -align -qopt-report=3 -qopt-report-phase=vec
micnativeloadex try
```



Oregon State University
Computer Graphics

mjb - March 30, 2015

Gaining Access to the Vector Units

```
#pragma omp simd
for( int i = 0; i < N; i++ )
    C[i] = A[i] * B[i] ;
```

```
#pragma omp parallel for simd
for( int i = 0; i < N; i++ )
    C[i] = A[i] * B[i] ;
```

```
C[0:N] = A[0:N] * B[0:N] ;
```

```
icpc -mmic -o try try.cpp -O3 -m -openmp -align -qopt-report=3 -qopt-report-phase=vec
micnativeloadex try
```



Oregon State University
Computer Graphics

mjb - March 30, 2015

Turning Off All Vectorization

```
icpc -mmic -o try try.cpp -O3 -lm -openmp -no-vec
micnativeloadex try
```

The only reason I can think of to do this is when running benchmarks to compare vector vs. scalar array processing.

The Intel compiler does a *great* job of automatically vectorizing where it can. **Warning:** just because you didn't deliberately vectorize your code doesn't mean it didn't end up vectorized! Use the "-no-vec" flag instead.



Oregon State University
Computer Graphics

mjb - March 30, 2015

Vectorizing Conditionals

```
#pragma omp simd
for( int i = 0; i < N; i++ )
{
    if( D[i] == 0 )
        C[i] = A[i] * B[i] ;
    else
        C[i] = A[i] + B[i] ;
}
```

In my tests, this was 3-4x as fast as this.

```
#pragma omp simd
for( int i = 0; i < N; i++ )
{
    C[i] = ( D[i] == 0 ) ? A[i] * B[i] : A[i] + B[i] ;
}
```



Oregon State University
Computer Graphics

mjb - March 30, 2015

Reducing a Vector

```
float f = __sec_reduce_add( A[0:N] );
float f = __sec_reduce_mul( A[0:N] );
float f = __sec_reduce_max( A[0:N] );
float f = __sec_reduce_min( A[0:N] );

int i = __sec_reduce_max_ind( A[0:N] );
int i = __sec_reduce_min_ind( A[0:N] );

boolean b = __sec_reduce_all_zero( A[0:N] );
boolean b = __sec_reduce_all_nonzero( A[0:N] );
boolean b = __sec_reduce_any_zero( A[0:N] );
boolean b = __sec_reduce_any_nonzero( A[0:N] );
```

You must specify the array length. An argument of **A[:]** will throw a compiler error.



Oregon State University
Computer Graphics

mjb - March 30, 2015

Reducing a Vector

```
float sum = 0.;
for( int i = 0; i < n; i++ )
{
    sum += A[i];
}
```

In my tests, this was the same speed as this.

```
float sum = __sec_reduce_add( A[0:N] );
```



Oregon State University
Computer Graphics

mjb - March 30, 2015

Elemental Vector Functions

```
for( int i = 0; i < N; i++ )
{
    C[i] = A[i] * B[i];
}
```

In my tests, this was 3x as fast as this.

```
__declspec(vector)
float vmul( float x, float y )
{
    return x*y;
}

for( int i = 0; i < N; i++ )
{
    C[i] = vmul( A[i], B[i] );
}
```



Oregon State University
Computer Graphics

mjb - March 30, 2015

Offload Mode

```
float *A = new float[NUMS];
float *B = new float[NUMS];
float *C = new float[NUMS];
```

```
...
double *dtp = new double;
double Time0 = omp_get_wtime( );
```

Data to send over

Data to bring back

```
#pragma offload target(micro) in(A:length(NUMS)) in(B:length(NUMS)) out(C:length(NUMS)) out(dtp:length(1))
```

```
{
    omp_set_num_threads( NUMT );
    double time0 = omp_get_wtime( );
```

```
#pragma omp parallel for simd
for( int i = 0; i < NUMS; i++ )
    C[i] = A[i] * B[i];
```

```
double time1 = omp_get_wtime( );
*dtp = time1 - time0;
```

```
}
```

```
double Time1 = omp_get_wtime( );
```

```
double overalldt = Time1 - Time0;
double offloaddt = *dtp;
```

```
fprintf( stderr, "%6d\t%6d\t%8.5f\t%8.5f\t%8.4f\t%8.2f\n", NUMT, NUMS,
    overalldt, offloaddt, 100.*offloaddt/overalldt, ((double)NUMS)/offloaddt/1000000. );
```

This
executes
on *rabbit*

This
executes
on the
Xeon Phi

This
executes
on *rabbit*



Oregon State University
Computer Graphics

mjb - March 30, 2015

Offload Mode

You don't need to do anything special with the compile line:

```
icpc -o try try.cpp -O3 -lm -openmp -align -qopt-report=3 -qopt-report-phase=vec
./try
```



Oregon State University
Computer Graphics

mjb - March 30, 2015

Offload Mode: Persistence Between Offloads

```
#define ALLOC  alloc_if(1)
#define REUSE  alloc_if(0)

#define RETAIN  free_if(0)
#define FREE   free_if(1)

#pragma offload target(mic) in(A:length(NUMS), ALLOC, RETAIN) out(C:length(NUMS), ALLOC, FREE )
{
    ...
}

...

#pragma offload target(mic) in(A:length(NUMS), REUSE, RETAIN) out(D:length(NUMS), ALLOC, RETAIN)
{
    ...
}

...

#pragma offload target(mic) in(A:length(NUMS), REUSE, FREE) out(D:length(NUMS), REUSE, FREE )
{
    ...
}
```



Alignment

To ensure alignment, replace this

```
float Temperature[NUMN];
```

with this:

```
#define ALIGN64    __declspec(align(64))
```

```
    . . .
```

```
ALIGN64 float Temperature[NUMN];
```



Oregon State University
Computer Graphics

mjb - March 30, 2015

Alignment

To ensure alignment, replace this

```
float *A = (float *) malloc( NUMS*sizeof(float) );  
float *B = (float *) malloc( NUMS*sizeof(float) );  
float *C = (float *) malloc( NUMS*sizeof(float) );
```

with this

```
float *A = (float *) _mm_malloc( NUMS*sizeof(float), 64 );  
float *B = (float *) _mm_malloc( NUMS*sizeof(float), 64 );  
float *C = (float *) _mm_malloc( NUMS*sizeof(float), 64 );
```

You then free memory with:

```
_mm_free( A );  
_mm_free( B );  
_mm_free( C );
```



Oregon State University
Computer Graphics

mjb - March 30, 2015

Alignment

If you want to ensure alignment, but still want to use C++'s *new* and *delete*, replace this:

```
float *A = new float [NUMS];
float *B = new float [NUMS];
float *C = new float [NUMS];
```

with this

```
float *pa = (float *) _mm_malloc( NUMS*sizeof(float), 64 );
float *pb = (float *) _mm_malloc( NUMS*sizeof(float), 64 );
float *pc = (float *) _mm_malloc( NUMS*sizeof(float), 64 );
```

```
float *A = new(pa) float [NUMS];
float *B = new(pb) float [NUMS];
float *C = new(pc) float [NUMS];
```

You then free memory with:

```
delete [ ] A;
delete [ ] B;
delete [ ] C;
```

An advantage of using *new* and *delete* instead of *malloc* is that they allow you to use C++ constructors and destructors.



Oregon State University
Computer Graphics

mjb - March 30, 2015

As You Create More and More Threads, On What Cores Do They End Up?

If you want them spread out onto as many cores as possible, execute this:

```
kmp_set_defaults( "KMP_AFFINITY=scatter" );
```

If you want them packed onto the first core until it has 4, then onto the second core until it has 4, etc., execute this:

```
kmp_set_defaults( "KMP_AFFINITY=compact" );
```



Use the scatter-mode if you want as much core-power applied to each thread as possible.

Use the compact-mode if there is an advantage to some threads sharing a core's local memory with other threads.



Oregon State University
Computer Graphics

mjb - March 30, 2015



Reservation System

<https://secure.engr.oregonstate.edu/engr/resources/bailey>

Bailey Resource Checkout Room Reservation System

11 Dec 2014 goto Help Search: Login

Areas: rabbit.engr

November 2014 December 2014 January 2015

Sun Mon Tue Wed Thu Fri Sat Sun Mon Tue Wed Thu Fri Sat Sun Mon Tue Wed Thu Fri Sat

2 3 4 5 6 7 8 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31

Thursday 11 December 2014

rabbit.engr

<<Go To Day Before Go To Today Go To Day After>>

Time:	rabbit.engr	Mike Bailey
07:00am	*	-
07:30am	*	-
08:00am	*	-
08:30am	*	-
09:00am	*	-
09:30am	*	-
10:00am	*	-
10:30am	*	-
11:00am	*	-
11:30am	*	-
12:00pm	*	-
12:30pm	*	-
01:00pm	*	-
01:30pm	*	-
02:00pm	*	-
02:30pm	*	-



Oregon State University
Computer Graphics

mjb - March 30, 2015

Getting you and your files over to the Xeon Phi:

Note: most of the time, this is unnecessary

```
ssh micuser@mic0
password: Intel
```

Much information is in:
/usr/local/apps/intel/studio.2015

```
scp file micuser@mic0:/home/micuser/file
```



Oregon State University
Computer Graphics

mjb - March 30, 2015